

PostgreSQL

Complete Mastery Guide

From Database Fundamentals to Advanced SQL

A Comprehensive Guide to PostgreSQL,
SQL Queries, Database Design, and Production-Grade
Database Management

Manoj S

`manojcs142006@gmail.com`

January 24, 2026

Contents

1 Introduction to Databases and PostgreSQL	7
1.1 What is a Database?	7
1.2 Why Do We Need Databases?	7
1.2.1 The Problem with File Storage	7
1.2.2 The Database Solution	7
1.3 Types of Databases	8
1.4 What is PostgreSQL?	8
1.4.1 Key Features of PostgreSQL	8
1.4.2 Why PostgreSQL?	8
1.5 SQL: Structured Query Language	8
1.5.1 SQL Categories	9
1.6 Database Concepts	9
1.6.1 Tables	9
1.6.2 Primary Key	9
1.6.3 Foreign Key	10
1.7 ACID Properties	10
1.8 Setting Up PostgreSQL	10
1.8.1 Installation Steps	10
1.8.2 Connecting to PostgreSQL	10
2 DDL: Data Definition Language	11
2.1 Creating Tables	11
2.1.1 Introduction	11
2.1.2 Syntax	11
2.1.3 Basic Table Creation	11
2.2 PostgreSQL Data Types	11
2.2.1 Numeric Types	11
2.2.2 Character Types	12
2.2.3 Date and Time Types	12
2.2.4 Boolean Type	12
2.3 Table Constraints	12
2.3.1 Introduction	12
2.3.2 Common Constraints	12
2.3.3 Advanced Table Creation	12
2.4 Viewing Table Structure	13
2.4.1 psql Commands	13
3 DML: Data Manipulation Language	14
3.1 INSERT: Adding Data	14
3.1.1 Introduction	14
3.1.2 Syntax	14
3.1.3 Single Row Insertion	14
3.1.4 Shorthand Insertion	14
3.1.5 Multiple Row Insertion	15
3.1.6 Inserting with DEFAULT Values	15
3.1.7 Constraint Violations	16
3.2 Sample Data Population	16
3.3 SELECT: Retrieving Data	16
3.3.1 Introduction	16

3.3.2	Basic Syntax	17
3.3.3	SELECT Examples	17
3.4	UPDATE: Modifying Data	18
3.4.1	Introduction	18
3.4.2	Syntax	18
3.4.3	UPDATE Examples	18
3.5	DELETE: Removing Data	19
3.5.1	Introduction	19
3.5.2	Syntax	19
3.5.3	DELETE Examples	19
3.6	TRUNCATE: Fast Delete	20
3.6.1	Introduction	20
3.6.2	Syntax	20
3.6.3	TRUNCATE vs DELETE	20
4	WHERE Clause: Filtering Data	22
4.1	Introduction	22
4.2	Syntax	22
4.3	Comparison Operators	22
4.4	Basic WHERE Examples	22
4.5	Logical Operators	23
4.5.1	AND Operator	23
4.5.2	OR Operator	23
4.6	IN Operator	24
4.6.1	Introduction	24
4.6.2	Syntax	24
4.7	NOT IN Operator	24
4.8	BETWEEN Operator	24
4.8.1	Introduction	24
4.8.2	Syntax	25
5	Sorting and Limiting Results	26
5.1	ORDER BY: Sorting Data	26
5.1.1	Introduction	26
5.1.2	Syntax	26
5.1.3	Sorting Direction	26
5.2	LIMIT: Restricting Results	27
5.2.1	Introduction	27
5.2.2	Syntax	27
5.3	OFFSET: Skipping Rows	27
5.3.1	Introduction	27
5.3.2	Syntax	28
5.4	Combining WHERE, ORDER BY, and LIMIT	28
6	DISTINCT and Pattern Matching	29
6.1	DISTINCT: Unique Values	29
6.1.1	Introduction	29
6.1.2	Syntax	29
6.2	LIKE: Pattern Matching	30
6.2.1	Introduction	30
6.2.2	Wildcards	30
6.2.3	Syntax	30

6.3	LIKE Examples	30
7	Aggregate Functions	32
7.1	Introduction	32
7.2	Common Aggregate Functions	32
7.3	COUNT Function	32
7.4	MAX and MIN Functions	32
7.5	SUM Function	33
7.6	AVG Function	33
8	GROUP BY: Grouping Data	35
8.1	Introduction	35
8.2	Syntax	35
8.3	Visualization	35
8.4	Basic GROUP BY Examples	35
8.5	GROUP BY with Different Aggregates	36
8.6	GROUP BY with Date Functions	37
9	HAVING Clause: Filtering Groups	39
9.1	Introduction	39
9.2	WHERE vs HAVING	39
9.3	Syntax	39
9.4	Query Execution Order	39
9.5	Basic HAVING Examples	40
9.6	WHERE + GROUP BY + HAVING	41
10	Interview Questions: Basic Level	42
10.1	Query Challenges	42
10.2	Advanced Filtering Questions	44
11	String Functions	46
11.1	Introduction	46
11.2	CONCAT: Concatenating Strings	46
11.2.1	Syntax	46
11.3	SUBSTRING: Extracting Substrings	47
11.3.1	Syntax	47
11.4	REPLACE: Replacing Substrings	48
11.4.1	Syntax	48
11.5	LENGTH: String Length	48
11.6	UPPER and LOWER: Case Conversion	49
11.7	LEFT and RIGHT: Edge Extraction	49
11.8	TRIM: Removing Whitespace	50
11.9	POSITION: Finding Substrings	51
11.10	String Function Practice Questions	51
12	Subqueries	53
12.1	Introduction	53
12.2	Types of Subqueries	53
12.3	Scalar Subqueries	53
12.4	Subquery Flow Diagram	54
12.5	Subqueries in Different Clauses	54

13 CASE Statements: Conditional Logic	56
13.1 Introduction	56
13.2 Syntax	56
13.3 Basic CASE Examples	56
13.4 CASE with Calculations	58
13.5 CASE with GROUP BY	59
13.6 Complex CASE Statements	59
13.7 CASE with Date Functions	61
14 DDL: Altering Tables	62
14.1 Introduction	62
14.2 Adding Columns	62
14.2.1 Syntax	62
14.3 Dropping Columns	62
14.4 Renaming Columns	63
14.5 Renaming Tables	63
14.6 Changing Data Types	63
14.7 Adding/Removing Defaults	64
14.8 Complete ALTER TABLE Examples	64
15 Advanced Constraints	65
15.1 CHECK Constraint	65
15.1.1 Introduction	65
15.1.2 Syntax	65
15.2 Testing Constraints	65
16 Interview Questions: Advanced Level	67
16.1 Filtering and Sorting Challenges	67
16.2 Aggregate Function Challenges	68
16.3 GROUP BY + HAVING Challenges	69
16.4 CASE Statement Challenges	70
16.5 Complex Subquery Challenges	71
17 Joins: Combining Tables	74
17.1 Introduction	74
17.2 Understanding Relationships	74
17.2.1 One-to-Many Relationship	74
17.3 Setting Up Tables for Joins	74
17.4 Relationship Diagram	74
17.5 Sample Data	75
17.6 Types of Joins	76
17.7 INNER JOIN	76
17.7.1 Syntax	76
17.8 LEFT JOIN (LEFT OUTER JOIN)	76
17.8.1 Introduction	77
17.9 JOIN with GROUP BY	77
17.10 JOIN with HAVING	78
17.11 Advanced JOIN Queries	78

18 Advanced JOIN Techniques	81
18.1 Subqueries in JOIN Context	81
18.2 Complex Filtering with JOIN	82
18.3 Advanced GROUP BY with JOIN	84
18.4 Multi-Condition Analysis	85
19 Multi-Table Joins: Complex Relationships	87
19.1 Introduction	87
19.2 E-Commerce Schema	87
19.3 Schema Diagram	88
19.4 Sample Data	88
19.5 Basic Multi-Table Queries	89
19.6 Aggregation Across Multiple Tables	90
19.7 Revenue Calculations	91
19.8 Advanced Multi-Table Queries	92
20 Views: Virtual Tables	96
20.1 Introduction	96
20.2 Why Use Views?	96
20.3 Creating Views	96
20.3.1 Syntax	96
20.4 Querying Views	97
20.5 Aggregation on Views	97
20.6 ROLLUP: Grand Totals	98
20.6.1 Introduction	98
20.6.2 Syntax	98
20.7 Managing Views	99
20.7.1 Dropping Views	99
20.7.2 Listing Views	99
21 Stored Procedures and Functions	100
21.1 Introduction	100
21.2 Stored Procedures	100
21.2.1 What is a Stored Procedure?	100
21.2.2 Syntax	100
21.3 User-Defined Functions (UDF)	101
21.3.1 What is a UDF?	101
21.3.2 Syntax	101
21.4 Procedures vs Functions	102
22 Window Functions: Advanced Analytics	103
22.1 Introduction	103
22.2 Window Function Syntax	103
22.3 Running Totals and Averages	103
22.4 PARTITION BY: Department-wise Analysis	104
22.5 Ranking Functions	105
22.6 Cumulative Sum per Department	106
23 LEAD and LAG Functions	108
23.1 Introduction	108
23.2 LAG Function	108
23.3 LEAD Function	109

24 CTEs: Common Table Expressions	110
24.1 Introduction	110
24.2 Syntax	110
24.3 CTE with GROUP BY	110
24.4 CTE with Window Functions	111
24.5 Multiple CTEs	112
24.6 CTE vs Subquery	113
24.7 Advanced CTE Examples	113
25 Triggers: Automated Database Actions	115
25.1 Introduction	115
25.2 Trigger Components	115
25.3 Creating a Trigger	115
25.3.1 Syntax	115
25.4 NEW and OLD Records	115
25.5 Example: Salary Validation Trigger	115
25.6 Removing Conflicting Constraints	116
25.7 Testing the Trigger	116
25.8 Trigger Types	117
25.9 Common Trigger Use Cases	117
25.10 Managing Triggers	117
25.10.1 Listing Triggers	118
25.10.2 Dropping Triggers	118
26 PostgreSQL Best Practices and Summary	119
26.1 Key Concepts Covered	119
26.2 SQL Best Practices	120
26.3 Query Optimization Tips	121
26.4 Common Pitfalls to Avoid	121
26.5 Essential PostgreSQL Commands	122
26.6 Next Steps for Learning	122

1 Introduction to Databases and PostgreSQL

1.1 What is a Database?

Definition

A database is an organized collection of structured data stored electronically in a computer system, designed for efficient storage, retrieval, and management of information.

Databases are the backbone of modern applications. They power:

- **Web Applications:** Facebook, Instagram, Amazon
- **Mobile Apps:** WhatsApp, Banking apps
- **Enterprise Systems:** HR, CRM, ERP systems
- **Data Analytics:** Business intelligence, ML pipelines

1.2 Why Do We Need Databases?

1.2.1 The Problem with File Storage

Before databases, data was stored in files. This approach had severe limitations:

Warning

Problems with File-Based Storage:

- **Data Redundancy:** Same data stored in multiple files
- **Data Inconsistency:** Updates in one file don't reflect in others
- **Difficult Data Access:** No standard way to query data
- **Security Issues:** Hard to implement access control
- **Concurrent Access:** Multiple users can't access simultaneously
- **No Data Integrity:** No way to enforce rules on data

1.2.2 The Database Solution

Databases solve all these problems by providing:

1. **Centralized Data Storage:** Single source of truth
2. **Data Integrity:** Enforce rules and constraints
3. **Concurrent Access:** Multiple users can work simultaneously
4. **Security:** Role-based access control
5. **Standardized Query Language:** SQL for all operations
6. **ACID Properties:** Ensures reliable transactions

1.3 Types of Databases

Type	Description	Examples
Relational	Data stored in tables with relationships	PostgreSQL, MySQL, Oracle
NoSQL	Flexible schema, document-based	MongoDB, Cassandra
Graph	Optimized for relationships	Neo4j, Amazon Neptune
Key-Value	Simple key-value pairs	Redis, DynamoDB

1.4 What is PostgreSQL?

PostgreSQL Definition

PostgreSQL is a powerful, open-source relational database management system (RDBMS) known for its robustness, scalability, and standards compliance.

1.4.1 Key Features of PostgreSQL

1. **Open Source:** Completely free to use
2. **ACID Compliant:** Ensures reliable transactions
3. **Advanced Data Types:** JSON, Arrays, Custom types
4. **Extensible:** Custom functions, operators, data types
5. **Concurrency:** Multi-Version Concurrency Control (MVCC)
6. **Performance:** Advanced indexing, query optimization
7. **Security:** Row-level security, SSL support
8. **Standards Compliant:** Follows SQL standards

1.4.2 Why PostgreSQL?

Advantage	Description
Community	Large, active community with excellent documentation
Enterprise-Ready	Used by major companies (Apple, Instagram, Spotify)
Feature-Rich	Advanced features like CTEs, Window Functions
Performance	Excellent query optimization and indexing
Reliability	Known for data integrity and crash recovery

1.5 SQL: Structured Query Language

What is SQL?

SQL (Structured Query Language) is a standardized programming language used to manage and manipulate relational databases.

1.5.1 SQL Categories

SQL commands are divided into five categories:

Category	Full Form	Purpose
DDL	Data Definition Language	Define/modify database structure CREATE, ALTER, DROP, TRUNCATE
DML	Data Manipulation Language	Manipulate data in tables INSERT, UPDATE, DELETE
DQL	Data Query Language	Retrieve data from database SELECT
DCL	Data Control Language	Control access permissions GRANT, REVOKE
TCL	Transaction Control Language	Manage transactions COMMIT, ROLLBACK, SAVEPOINT

1.6 Database Concepts

1.6.1 Tables

Table Definition

A **table** is a collection of related data organized in rows and columns, similar to a spreadsheet.

Components:

- **Rows (Records/Tuples)**: Individual data entries
- **Columns (Fields/Attributes)**: Data categories

1.6.2 Primary Key

Primary Key

A **Primary Key** is a column (or set of columns) that uniquely identifies each row in a table.

Characteristics:

- Must be unique
- Cannot be NULL
- Only one primary key per table
- Often auto-incrementing

1.6.3 Foreign Key

Foreign Key

A **Foreign Key** is a column that creates a link between two tables by referencing the primary key of another table.

Purpose: Maintains referential integrity between related tables

1.7 ACID Properties

ACID in Databases

ACID properties ensure reliable database transactions:

1. **Atomicity:** Transaction is all-or-nothing
2. **Consistency:** Database remains in valid state
3. **Isolation:** Concurrent transactions don't interfere
4. **Durability:** Committed changes are permanent

1.8 Setting Up PostgreSQL

1.8.1 Installation Steps

1. Download PostgreSQL from official website
2. Run installer for your operating system
3. Set password for default user (postgres)
4. Note the port number (default: 5432)
5. Install pgAdmin (GUI tool) if desired

1.8.2 Connecting to PostgreSQL

Using Command Line:

```
# Windows  
psql -U postgres  
  
# Linux/Mac  
sudo -u postgres psql
```

Using pgAdmin:

- Open pgAdmin
- Connect to server
- Enter password
- Navigate to databases

2 DDL: Data Definition Language

2.1 Creating Tables

2.1.1 Introduction

Creating tables is the foundation of database design. A table structure defines:

- What data will be stored
- What types of data are allowed
- What constraints apply to the data

2.1.2 Syntax

Syntax

```

1 CREATE TABLE table_name (
2     column1 datatype [constraints],
3     column2 datatype [constraints],
4     ...
5     [table_constraints]
6 );

```

2.1.3 Basic Table Creation

Simple Students Table

Requirement: Create a table to store student information

```

1 CREATE TABLE students(
2     id INT,
3     name VARCHAR(50),
4     city VARCHAR(50)
5 );

```

Listing 1: Creating Students Table

Explanation:

- **id:** Integer type for student ID
- **name:** Variable character string up to 50 characters
- **city:** Variable character string up to 50 characters

2.2 PostgreSQL Data Types

2.2.1 Numeric Types

Type	Description	Example
INT	Integer (4 bytes)	42, -100, 0
SERIAL	Auto-incrementing integer	1, 2, 3...
NUMERIC(p,s)	Exact decimal	NUMERIC(10,2) for currency
FLOAT	Floating-point number	3.14159

2.2.2 Character Types

Type	Description
VARCHAR(n)	Variable-length string up to n characters
CHAR(n)	Fixed-length string, always n characters
TEXT	Unlimited variable-length string

2.2.3 Date and Time Types

Type	Format	Example
DATE	YYYY-MM-DD	2024-01-15
TIME	HH:MI:SS	14:30:00
TIMESTAMP	Date + Time	2024-01-15 14:30:00

2.2.4 Boolean Type

- BOOLEAN: Stores TRUE, FALSE, or NULL

2.3 Table Constraints

2.3.1 Introduction

Constraints enforce rules on data to maintain integrity and validity.

2.3.2 Common Constraints

Constraint	Purpose
PRIMARY KEY	Uniquely identifies each row
FOREIGN KEY	Links to another table's primary key
NOT NULL	Column cannot contain NULL values
UNIQUE	All values must be different
CHECK	Values must satisfy a condition
DEFAULT	Provides default value if none specified

2.3.3 Advanced Table Creation

Employees Table with Constraints

Requirement: Create a production-ready employees table

```

1 CREATE TABLE employees(
2     emp_id SERIAL PRIMARY KEY,
3     fname VARCHAR(50) NOT NULL,
4     lname VARCHAR(50) NOT NULL,
5     email VARCHAR(100) NOT NULL UNIQUE,
6     dept VARCHAR(20),
7     salary NUMERIC(10,2) DEFAULT 30000.00,
8     hire_date DATE NOT NULL DEFAULT CURRENT_DATE
9 );

```

Listing 2: Employees Table with All Constraints

Constraint Breakdown:

- SERIAL PRIMARY KEY: Auto-incrementing unique identifier

- NOT NULL: First name, last name, email, hire date required
- UNIQUE: Email must be unique across all employees
- DEFAULT 30000.00: Salary defaults to 30,000 if not specified
- DEFAULT CURRENT_DATE: Hire date defaults to today

Important Note

SERIAL Data Type:

SERIAL is PostgreSQL-specific and automatically:

- Creates a sequence
- Auto-increments with each new row
- Commonly used for primary keys
- Equivalent to AUTO_INCREMENT in MySQL

2.4 Viewing Table Structure

2.4.1 psql Commands

```
-- List all tables
\dt

-- Describe table structure
\d table_name

-- Detailed table info
\dt+ table_name
```

3 DML: Data Manipulation Language

3.1 INSERT: Adding Data

3.1.1 Introduction

The INSERT statement adds new rows to a table. This is how you populate your database with data.

3.1.2 Syntax

Syntax

```

1 -- Single row insertion
2 INSERT INTO table_name (column1, column2, ...)
3 VALUES (value1, value2, ...);
4
5 -- Multiple rows insertion
6 INSERT INTO table_name (column1, column2, ...)
7 VALUES
8     (value1, value2, ...),
9     (value1, value2, ...),
10    (value1, value2, ...);
11
12 -- Insert all columns (order must match table definition)
13 INSERT INTO table_name
14 VALUES (value1, value2, ...);
```

3.1.3 Single Row Insertion

Inserting One Student

```

1 INSERT INTO students(id, name, city)
2 VALUES (101, 'Manoj', 'bangalore');
```

Listing 3: Single Row INSERT

Result:

- Adds one row to students table
- id = 101, name = 'Manoj', city = 'bangalore'

3.1.4 Shorthand Insertion

Omitting Column Names

```

1 INSERT INTO students
2 VALUES (102, 'Manu', 'Mysore');
```

Listing 4: Insert Without Column Names

Note: Values must be in the exact order as table columns

Warning

Best Practice: Always specify column names explicitly!

Why?

- More readable and maintainable
- Prevents errors if table structure changes
- Self-documenting code
- Allows inserting columns in any order

3.1.5 Multiple Row Insertion

Bulk Insert

```

1 INSERT INTO students VALUES
2 (103, 'sham', 'Mumbai'),
3 (104, 'Akash', 'Gujarat'),
4 (105, 'chinmai', 'mandya');
```

Listing 5: Multiple Rows at Once

Advantages:

- More efficient than individual inserts
- Single transaction
- Faster execution

3.1.6 Inserting with DEFAULT Values

Using DEFAULT Constraints

```

1 -- Insert with some default values
2 INSERT INTO employees (fname, lname, email, dept, salary)
3 VALUES
4 ('Manoj', 'S', 'manoj.s@company.com', 'IT', 45000.00),
5 ('Akash', 'Kumar', 'akash.k@company.com', 'HR', 38000.00);
6
7 -- Insert using all defaults for salary and hire_date
8 INSERT INTO employees(fname, lname, email, dept)
9 VALUES ('Rohit', 'Shetty', 'rohitshetty7@company.com', 'IT');
```

Listing 6: Insert Using Defaults

What happens:

- `emp_id`: Auto-generated (SERIAL)
- `salary`: Uses default 30000.00
- `hire_date`: Uses current date

3.1.7 Constraint Violations

UNIQUE Constraint Error

```

1 -- This will FAIL - duplicate email
2 INSERT INTO employees(fname, lname, email, dept)
3 VALUES ('Manoj', 'Shetty', 'rohitshetty7@company.com', 'IT');

```

Listing 7: Attempting Duplicate Email

Error Message:

```

ERROR: duplicate key value violates unique constraint
      "employees_email_key"
DETAIL: Key (email)=(rohitshetty7@company.com) already exists.

```

Reason: Email 'rohitshetty7@company.com' already exists in table

3.2 Sample Data Population

```

1 -- Populate employees table with sample data
2 INSERT INTO employees (fname, lname, email, dept, salary, hire_date)
3 VALUES
4 ('Raj', 'Sharma', 'raj.sharma@example.com', 'IT', 50000.00, '2020-01-15'),
5 ('Priya', 'Singh', 'priya.singh@example.com', 'HR', 45000.00, '2019-03-22'),
6 ('Arjun', 'Verma', 'arjun.verma@example.com', 'IT', 55000.00, '2021-06-01'),
7 ('Suman', 'Patel', 'suman.patel@example.com', 'Finance', 60000.00, '2018-07-30'),
8 ('Kavita', 'Rao', 'kavita.rao@example.com', 'HR', 47000.00, '2020-11-10'),
9 ('Amit', 'Gupta', 'amit.gupta@example.com', 'Marketing', 52000.00, '2020-09-25'),
10 ('Neha', 'Desai', 'neha.desai@example.com', 'IT', 48000.00, '2019-05-18'),
11 ('Rahul', 'Kumar', 'rahul.kumar@example.com', 'IT', 53000.00, '2021-02-14'),
12 ('Anjali', 'Mehta', 'anjali.mehta@example.com', 'Finance', 61000.00, '2018-12-03'),
13 ('Vijay', 'Nair', 'vijay.nair@example.com', 'Marketing', 50000.00, '2020-04-19');

```

Listing 8: Comprehensive Employee Data

3.3 SELECT: Retrieving Data

3.3.1 Introduction

The SELECT statement is the most commonly used SQL command. It retrieves data from one or more tables.

3.3.2 Basic Syntax

Syntax

```

1 SELECT column1, column2, ...
2 FROM table_name;
3
4 -- Select all columns
5 SELECT * FROM table_name;

```

3.3.3 SELECT Examples

Basic SELECT Queries

```

1 -- Select all columns, all rows
2 SELECT * FROM students;
3
4 -- Select specific column
5 SELECT id FROM students;
6
7 -- Select multiple columns
8 SELECT id, name FROM students;

```

Listing 9: Various SELECT Statements

Output for `SELECT * FROM students`:

Output

id	name	city
101	Manoj	bangalore
102	Manu	Mysore
103	sham	Mumbai
104	Akash	Gujarat
105	chinmai	mandya

(5 rows)

Output for `SELECT id, name FROM students`:

Output

id	name
101	Manoj
102	Manu
103	sham
104	Akash
105	chinmai

(5 rows)

Important Note

Performance Tip:

Always select only the columns you need!

- `SELECT *` retrieves all columns (slower, more data)
- Specify exact columns for better performance
- Especially important for tables with many columns

3.4 UPDATE: Modifying Data

3.4.1 Introduction

The `UPDATE` statement modifies existing data in a table. Use it carefully - without a `WHERE` clause, it updates ALL rows!

3.4.2 Syntax

Syntax

```
1 UPDATE table_name
2 SET column1 = value1, column2 = value2, ...
3 WHERE condition;
```

3.4.3 UPDATE Examples

Updating Student City

```
1 UPDATE students
2 SET city = 'Pune'
3 WHERE id = 104;
```

Listing 10: `UPDATE` with `WHERE` Clause

Before UPDATE:

<code>id</code>	<code>name</code>	<code>city</code>
104	Akash	Gujarat

After UPDATE:

<code>id</code>	<code>name</code>	<code>city</code>
104	Akash	Pune

Explanation

- Finds row where `id = 104`
- Changes city from 'Gujarat' to 'Pune'

- Only one row affected

Warning

CRITICAL WARNING:

ALWAYS use WHERE clause with UPDATE!

```
1 -- This updates ALL rows!
2 UPDATE students
3 SET city = 'Pune'; -- DANGER: Updates all students!
```

Listing 11: Dangerous UPDATE Without WHERE

Consequence: Every student's city becomes 'Pune'

3.5 DELETE: Removing Data

3.5.1 Introduction

The DELETE statement removes rows from a table. Like UPDATE, it requires careful use of the WHERE clause.

3.5.2 Syntax

Syntax

```
1 DELETE FROM table_name
2 WHERE condition;
```

3.5.3 DELETE Examples

Deleting Specific Student

```
1 DELETE FROM students
2 WHERE id = 105;
```

Listing 12: DELETE with WHERE Clause

Before DELETE:

id	name	city
101	Manoj	bangalore
102	Manu	Mysore
103	sham	Mumbai
104	Akash	Pune
105	chinmai	mandya

(5 rows)

After DELETE:

id	name	city
101	Manoj	bangalore
102	Manu	Mysore
103	sham	Mumbai
104	Akash	Pune

```

101 | Manoj    | bangalore
102 | Manu     | Mysore
103 | sham     | Mumbai
104 | Akash    | Pune
(4 rows)

```

Result: Row with id=105 permanently removed

Warning

CRITICAL WARNING:

ALWAYS use WHERE clause with DELETE!

```

1 -- This deletes ALL rows!
2 DELETE FROM students; -- DANGER: Removes all students!

```

Listing 13: Dangerous DELETE Without WHERE

Consequence: Entire table emptied (but structure remains)

3.6 TRUNCATE: Fast Delete

3.6.1 Introduction

TRUNCATE removes all rows from a table quickly. It's faster than DELETE but less flexible.

3.6.2 Syntax

Syntax

```
1 TRUNCATE TABLE table_name;
```

3.6.3 TRUNCATE vs DELETE

Aspect	TRUNCATE	DELETE
Speed	Very fast	Slower
WHERE Clause	Not supported	Supported
Rollback	Cannot rollback	Can rollback
Triggers	Doesn't fire triggers	Fires triggers
Auto-increment	Resets to 1	Doesn't reset
Use Case	Remove all data	Remove specific rows

TRUNCATE Example

```

1 -- Remove all employees (fast)
2 TRUNCATE TABLE employees;

```

Listing 14: Using TRUNCATE

Result

- All rows deleted instantly
- Table structure remains
- SERIAL counter resets to 1
- Much faster than DELETE for large tables

4 WHERE Clause: Filtering Data

4.1 Introduction

The WHERE clause filters rows based on conditions. It's used with SELECT, UPDATE, and DELETE to target specific data.

4.2 Syntax

Syntax

```
1 SELECT column1, column2, ...
2 FROM table_name
3 WHERE condition;
```

4.3 Comparison Operators

Operator	Meaning	Example
=	Equal to	WHERE dept = 'IT'
<> or !=	Not equal to	WHERE dept <> 'HR'
>	Greater than	WHERE salary > 50000
<	Less than	WHERE age < 30
>=	Greater or equal	WHERE salary >= 50000
<=	Less or equal	WHERE age <= 25

4.4 Basic WHERE Examples

Simple WHERE Conditions

```
1 -- Employees in IT department
2 SELECT * FROM employees WHERE dept = 'IT';
3
4 -- Employee named Raj
5 SELECT * FROM employees WHERE fname = 'Raj';
6
7 -- Employee with specific ID
8 SELECT * FROM employees WHERE emp_id = 5;
9
10 -- Employees earning 50,000 or more
11 SELECT * FROM employees WHERE salary >= 50000;
```

Listing 15: Basic Filtering

Output for WHERE dept = 'IT':

Output

emp_id	fname	lname	email	dept	salary
1	Raj	Sharma	raj.sharma@...	IT	50000.00
3	Arjun	Verma	arjun.verma@...	IT	55000.00
7	Neha	Desai	neha.desai@...	IT	48000.00
8	Rahul	Kumar	rahul.kumar@...	IT	53000.00

(4 rows)

4.5 Logical Operators

4.5.1 AND Operator

AND Operator

Returns rows where **ALL** conditions are true

Using AND

```

1 -- IT employees earning 50,000 or less
2 SELECT * FROM employees
3 WHERE dept = 'IT' AND salary <= 50000;

```

Listing 16: Multiple Conditions with AND

Output:

Output

emp_id	fname	lname	email	dept	salary
1	Raj	Sharma	raj.sharma@...	IT	50000.00
7	Neha	Desai	neha.desai@...	IT	48000.00

(2 rows)

Logic: Both conditions must be satisfied

4.5.2 OR Operator

OR Operator

Returns rows where **ANY** condition is true

Using OR

```

1 -- Employees in IT OR HR departments
2 SELECT * FROM employees
3 WHERE dept = 'IT' OR dept = 'HR';

```

Listing 17: Multiple Conditions with OR

Result: Returns employees from both IT and HR departments

4.6 IN Operator

4.6.1 Introduction

The IN operator is a cleaner alternative to multiple OR conditions.

4.6.2 Syntax

Syntax

```
1 SELECT * FROM table_name
2 WHERE column IN (value1, value2, value3, ...);
```

IN Operator Examples

```
1 -- Employees in HR or IT (recommended way)
2 SELECT * FROM employees
3 WHERE dept IN ('HR', 'IT');
4
5 -- Equivalent to:
6 SELECT * FROM employees
7 WHERE dept = 'HR' OR dept = 'IT';
```

Listing 18: Using IN Operator

Advantages of IN:

- More readable
- Easier to maintain
- Works with subqueries
- Shorter syntax

4.7 NOT IN Operator

Excluding Values

```
1 -- Employees NOT in HR or IT departments
2 SELECT * FROM employees
3 WHERE dept NOT IN ('HR', 'IT');
```

Listing 19: Using NOT IN

Result: Returns employees from Finance and Marketing only

4.8 BETWEEN Operator

4.8.1 Introduction

BETWEEN selects values within a range (inclusive of both endpoints).

4.8.2 Syntax

Syntax

```
1 SELECT * FROM table_name  
2 WHERE column BETWEEN value1 AND value2;
```

BETWEEN Examples

```
1 -- Employees earning between 40,000 and 60,000 (inclusive)  
2 SELECT * FROM employees  
3 WHERE salary BETWEEN 40000 AND 60000;  
4  
5 -- Equivalent to:  
6 SELECT * FROM employees  
7 WHERE salary >= 40000 AND salary <= 60000;
```

Listing 20: Using BETWEEN

Output:

Output

fname		salary
Raj		50000.00
Priya		45000.00
Arjun		55000.00
Suman		60000.00
Kavita		47000.00
Amit		52000.00
Neha		48000.00
Rahul		53000.00
Vijay		50000.00
(9 rows)		

Important Note

BETWEEN is Inclusive:

BETWEEN 40000 AND 60000 includes both 40,000 and 60,000.

5 Sorting and Limiting Results

5.1 ORDER BY: Sorting Data

5.1.1 Introduction

ORDER BY sorts query results based on one or more columns.

5.1.2 Syntax

Syntax

```
1 SELECT column1, column2, ...
2 FROM table_name
3 ORDER BY column1 [ASC|DESC], column2 [ASC|DESC];
```

5.1.3 Sorting Direction

- ASC: Ascending order (default)
- DESC: Descending order

ORDER BY Examples

```
1 -- Sort by salary (ascending - default)
2 SELECT * FROM employees ORDER BY salary;
3
4 -- Sort by department (descending)
5 SELECT * FROM employees ORDER BY dept DESC;
6
7 -- Sort by multiple columns
8 SELECT * FROM employees
9 ORDER BY dept ASC, salary DESC;
```

Listing 21: Sorting Examples

Output for ORDER BY salary:

Output

fname	dept	salary
Priya	HR	45000.00
Kavita	HR	47000.00
Neha	IT	48000.00
Raj	IT	50000.00
Vijay	Marketing	50000.00
Amit	Marketing	52000.00
Rahul	IT	53000.00
Arjun	IT	55000.00
Suman	Finance	60000.00
Anjali	Finance	61000.00
(10 rows)		

Multi-column sort explanation:

- First sorts by department (A-Z)
- Within each department, sorts by salary (high to low)

5.2 LIMIT: Restricting Results

5.2.1 Introduction

LIMIT restricts the number of rows returned by a query.

5.2.2 Syntax

Syntax

```
1 SELECT * FROM table_name
2 LIMIT number_of_rows;
```

LIMIT Examples

```
1 -- Get first 5 employees
2 SELECT * FROM employees LIMIT 5;
3
4 -- Get top 3 highest paid employees
5 SELECT * FROM employees
6 ORDER BY salary DESC
7 LIMIT 3;
```

Listing 22: Using LIMIT

Output for top 3 salaries:

Output

fname	dept	salary
Anjali	Finance	61000.00
Suman	Finance	60000.00
Arjun	IT	55000.00

(3 rows)

5.3 OFFSET: Skipping Rows

5.3.1 Introduction

OFFSET skips a specified number of rows before returning results. Often used with LIMIT for pagination.

5.3.2 Syntax

Syntax

```
1 SELECT * FROM table_name
2 LIMIT number_of_rows OFFSET number_to_skip;
```

Pagination Example

```
1 -- Page 1: First 5 employees
2 SELECT * FROM employees
3 LIMIT 5 OFFSET 0;
4
5 -- Page 2: Next 5 employees
6 SELECT * FROM employees
7 LIMIT 5 OFFSET 5;
8
9 -- Page 3: Next 5 employees
10 SELECT * FROM employees
11 LIMIT 5 OFFSET 10;
```

Listing 23: Implementing Pagination

Pagination Formula:

- Page size = LIMIT value
- Skip amount = (page_number - 1) × page_size

5.4 Combining WHERE, ORDER BY, and LIMIT

Complex Query Example

```
1 -- Top 7 employees earning between 30k-65k, sorted by salary
2 SELECT * FROM employees
3 WHERE salary BETWEEN 30000 AND 65000
4 ORDER BY salary DESC
5 LIMIT 7 OFFSET 5;
```

Listing 24: Complete Query with Multiple Clauses

Execution Order:

1. WHERE: Filter employees by salary range
2. ORDER BY: Sort filtered results by salary
3. LIMIT OFFSET: Skip 5, return next 7

6 DISTINCT and Pattern Matching

6.1 DISTINCT: Unique Values

6.1.1 Introduction

DISTINCT removes duplicate values from query results, showing only unique values.

6.1.2 Syntax

Syntax

```
1 SELECT DISTINCT column1, column2, ...
2 FROM table_name;
```

DISTINCT Examples

```
1 -- Get all unique departments
2 SELECT DISTINCT dept FROM employees;
```

Listing 25: Finding Unique Values

Output:

Output

```
dept
-----
IT
HR
Finance
Marketing
(4 rows)
```

Without DISTINCT:

```
dept
-----
IT
HR
IT
Finance
HR
Marketing
IT
IT
Finance
Marketing
(10 rows)
```

Use Cases:

- Finding unique categories
- Counting distinct values

- Removing duplicates from results

6.2 LIKE: Pattern Matching

6.2.1 Introduction

LIKE operator performs pattern matching on text using wildcards.

6.2.2 Wildcards

Wildcard	Meaning
%	Matches zero or more characters
_	Matches exactly one character

6.2.3 Syntax

Syntax

```
1 SELECT * FROM table_name
2 WHERE column LIKE 'pattern';
```

6.3 LIKE Examples

Pattern Matching Scenarios

```
1 -- Names starting with 'A'
2 SELECT * FROM employees WHERE fname LIKE 'A%';
3
4 -- Names ending with 'a'
5 SELECT * FROM employees WHERE fname LIKE '%a';
6
7 -- Names containing 'i'
8 SELECT * FROM employees WHERE fname LIKE '%i%';
9
10 -- 'i' as second letter
11 SELECT * FROM employees WHERE fname LIKE '_i%';
12
13 -- Names containing 'john' (case-insensitive in PostgreSQL)
14 SELECT * FROM employees WHERE fname LIKE '%john%';
15
16 -- Department with exactly 2 characters
17 SELECT * FROM employees WHERE dept LIKE '__';
```

Listing 26: Various LIKE Patterns

Output for fname LIKE 'A%':

Output

```
fname | lname | dept
-----+-----+-----
Arjun | Verma | IT
Amit  | Gupta | Marketing
Anjali | Mehta | Finance
(3 rows)
```

Output for fname LIKE '%i%':

Output

```
fname
-----
Priya
Kavita
Amit
Vijay
(4 rows)
```

Important Note**Case Sensitivity:**

- LIKE is case-sensitive in many databases
- In PostgreSQL, use ILIKE for case-insensitive matching
- Example: WHERE fname ILIKE 'a%' matches 'Amit' and 'amit'

7 Aggregate Functions

7.1 Introduction

Aggregate functions perform calculations on multiple rows and return a single value.

7.2 Common Aggregate Functions

Function	Purpose
COUNT()	Counts number of rows
SUM()	Calculates total sum
AVG()	Calculates average
MAX()	Finds maximum value
MIN()	Finds minimum value

7.3 COUNT Function

Counting Rows

```

1 -- Total number of employees
2 SELECT COUNT(*) FROM employees;
3
4 -- Count employees in IT department
5 SELECT COUNT(*) FROM employees WHERE dept = 'IT';

```

Listing 27: COUNT Examples

Output:

Output

```

-- COUNT(*)
count
-----
10
(1 row)

-- COUNT with WHERE
count
-----
4
(1 row)

```

7.4 MAX and MIN Functions

Finding Extremes

```

1 -- Highest salary
2 SELECT MAX(salary) FROM employees;
3
4 -- Lowest salary
5 SELECT MIN(salary) FROM employees;
6

```

```

7 -- Highest salary in IT department
8 SELECT MAX(salary) FROM employees WHERE dept = 'IT';

```

Listing 28: MAX and MIN Examples

Output:

Output

```

-- MAX(salary)
max
-----
61000.00
(1 row)

-- MIN(salary)
min
-----
45000.00
(1 row)

```

7.5 SUM Function

Calculating Totals

```

1 -- Total salary expense
2 SELECT SUM(salary) FROM employees;
3
4 -- Total salary for IT department
5 SELECT SUM(salary) FROM employees WHERE dept = 'IT';

```

Listing 29: SUM Examples

Output:

Output

```

-- SUM(salary)
sum
-----
521000.00
(1 row)

```

7.6 AVG Function

Calculating Averages

```

1 -- Average salary across company
2 SELECT AVG(salary) FROM employees;
3
4 -- Average salary for employees under 30
5 SELECT AVG(salary) FROM employees WHERE age < 30;

```

```
6  
7 -- Rounded average  
8 SELECT ROUND(AVG(salary), 2) AS avg_salary  
9 FROM employees;
```

Listing 30: AVG Examples

Output:**Output**

```
-- AVG(salary)  
avg  
-----  
52100.0000000000000000  
(1 row)  
  
-- Rounded AVG  
avg_salary  
-----  
52100.00  
(1 row)
```

8 GROUP BY: Grouping Data

8.1 Introduction

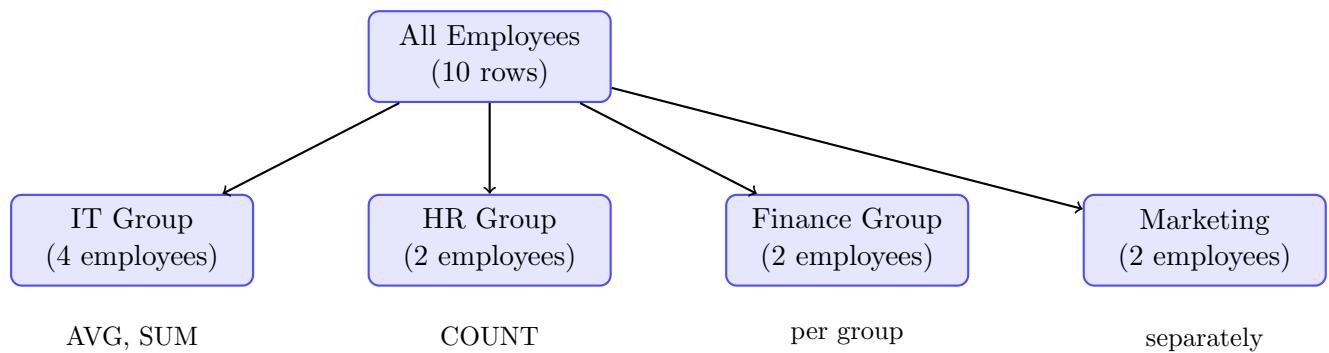
GROUP BY groups rows with the same values in specified columns, allowing aggregate functions to be applied to each group.

8.2 Syntax

Syntax

```
1 SELECT column1, aggregate_function(column2)
2 FROM table_name
3 GROUP BY column1;
```

8.3 Visualization



8.4 Basic GROUP BY Examples

Counting Employees per Department

```
1 -- Employees per department
2 SELECT dept, COUNT(*) AS emp_count
3 FROM employees
4 GROUP BY dept;
```

Listing 31: COUNT with GROUP BY

Output:

Output

dept	emp_count
IT	4
HR	2
Finance	2
Marketing	2
(4 rows)	

Explanation:

- Groups all employees by department
- Counts employees in each group
- Returns one row per department

8.5 GROUP BY with Different Aggregates

Multiple Aggregate Examples

```

1 -- Average salary per department
2 SELECT dept, AVG(salary) AS avg_salary
3 FROM employees
4 GROUP BY dept;
5
6 -- Maximum salary per department
7 SELECT dept, MAX(salary) AS max_salary
8 FROM employees
9 GROUP BY dept;
10
11 -- Minimum salary per department
12 SELECT dept, MIN(salary) AS min_salary
13 FROM employees
14 GROUP BY dept;
15
16 -- Total salary per department
17 SELECT dept, SUM(salary) AS total_salary
18 FROM employees
19 GROUP BY dept;

```

Listing 32: Various Aggregates by Department

Output for AVG(salary):

Output

dept	avg_salary
IT	51500.0000000000000000
HR	46000.0000000000000000
Finance	60500.0000000000000000
Marketing	51000.0000000000000000
(4 rows)	

Output for SUM(salary):

Output

dept	total_salary
IT	206000.00
HR	92000.00
Finance	121000.00
Marketing	102000.00

(4 rows)

8.6 GROUP BY with Date Functions

Grouping by Year

```

1 -- Employees hired per year
2 SELECT EXTRACT(YEAR FROM hire_date) AS hire_year,
3       COUNT(*) AS emp_count
4 FROM employees
5 GROUP BY EXTRACT(YEAR FROM hire_date);

```

Listing 33: Employees Hired per Year

Output:**Output**

hire_year	emp_count
2018	2
2019	2
2020	4
2021	2

(4 rows)

Explanation:

- EXTRACT(YEAR FROM hire_date): Gets year from date
- Groups by hiring year
- Counts employees hired each year

Important Note**GROUP BY Rule:**

Any column in SELECT that's not in an aggregate function MUST be in GROUP BY clause.

```

1 -- Correct
2 SELECT dept, COUNT(*)
3 FROM employees
4 GROUP BY dept;
5

```

```
6 -- Wrong - fname not in GROUP BY
7 SELECT fname, dept, COUNT(*)
8 FROM employees
9 GROUP BY dept;
```

9 HAVING Clause: Filtering Groups

9.1 Introduction

HAVING filters groups created by GROUP BY, similar to how WHERE filters individual rows.

9.2 WHERE vs HAVING

Aspect	WHERE	HAVING
Filters	Individual rows	Groups
Used with	Any SELECT	GROUP BY
Timing	Before grouping	After grouping
Aggregates	Cannot use	Can use
Example	WHERE salary > 50000	HAVING AVG(salary) > 50000

9.3 Syntax

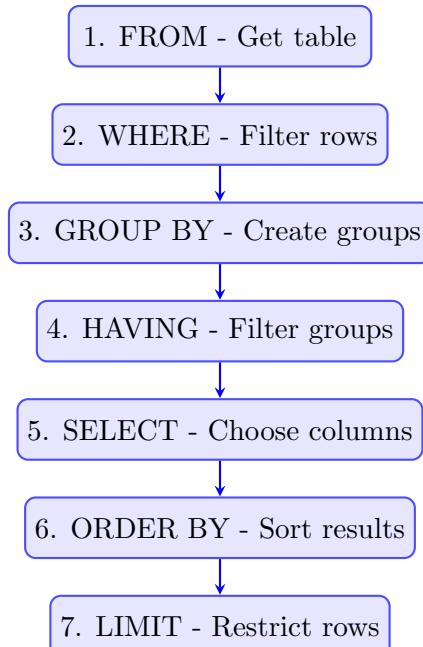
Syntax

```

1 SELECT column1, aggregate_function(column2)
2 FROM table_name
3 GROUP BY column1
4 HAVING condition;

```

9.4 Query Execution Order



9.5 Basic HAVING Examples

Filtering Groups

```

1 -- Departments with more than 2 employees
2 SELECT dept, COUNT(*) AS emp_count
3 FROM employees
4 GROUP BY dept
5 HAVING COUNT(*) > 2;

```

Listing 34: Departments with More Than 2 Employees

Output:

Output

dept	emp_count
IT	4

(1 row)

Explanation:

- Groups employees by department
- Counts employees per department
- Filters to show only departments with ≥ 2 employees
- Only IT has more than 2 employees

Departments with High Average Salary

```

1 -- Departments with average salary > 50000
2 SELECT dept, AVG(salary) AS avg_salary
3 FROM employees
4 GROUP BY dept
5 HAVING AVG(salary) > 50000;

```

Listing 35: Average Salary Filter

Output:

Output

dept	avg_salary
IT	51500.000000000000000000
Finance	60500.000000000000000000
Marketing	51000.000000000000000000

(3 rows)

Years with Multiple Hires

```

1 -- Years with more than 2 hires
2 SELECT EXTRACT(YEAR FROM hire_date) AS hire_year,
3       COUNT(*) AS emp_count
4 FROM employees
5 GROUP BY EXTRACT(YEAR FROM hire_date)
6 HAVING COUNT(*) > 2;

```

Listing 36: Filtering by Hire Count

Output:

Output

hire_year	emp_count
2020	4

(1 row)

9.6 WHERE + GROUP BY + HAVING

Complete Query Example

```

1 -- Departments (excluding HR) with more than 1 employee
2 SELECT dept, COUNT(*) AS emp_count
3 FROM employees
4 WHERE dept <> 'HR'
5 GROUP BY dept
6 HAVING COUNT(*) > 1;

```

Listing 37: Combining WHERE and HAVING

Execution Steps:

1. WHERE: Remove HR employees first
2. GROUP BY: Group remaining employees by department
3. HAVING: Keep only departments with > 1 employee

Output:

Output

dept	emp_count
IT	4
Finance	2
Marketing	2

(3 rows)

10 Interview Questions: Basic Level

10.1 Query Challenges

Question 1: Department Analysis

Q: Find the average salary by department for employees hired after 2019.

```
1 SELECT dept, AVG(salary) AS avg_salary
2 FROM employees
3 WHERE hire_date > '2019-12-31'
4 GROUP BY dept;
```

Listing 38: Solution

Output:

Output

dept	avg_salary
IT	52000.000000000000000000
HR	47000.000000000000000000
Marketing	51000.000000000000000000

(3 rows)

Question 2: Salary Range Filter

Q: Find departments where total salary is between 100,000 and 200,000.

```
1 SELECT dept, SUM(salary) AS total_salary
2 FROM employees
3 GROUP BY dept
4 HAVING SUM(salary) BETWEEN 100000 AND 200000;
```

Listing 39: Solution

Output:

Output

dept	total_salary
Finance	121000.00
Marketing	102000.00

(2 rows)

Question 3: Highest Salary Department

Q: Find the department with the highest salary.

```
1 SELECT dept, MAX(salary) AS max_salary
2 FROM employees
3 GROUP BY dept
4 ORDER BY max_salary DESC
5 LIMIT 1;
```

Listing 40: Solution

Output:**Output**

dept	max_salary
Finance	61000.00

(1 row)

Question 4: Best Average Salary**Q:** Find the department with the highest average salary.

```

1 SELECT dept, AVG(salary) AS avg_salary
2 FROM employees
3 GROUP BY dept
4 ORDER BY avg_salary DESC
5 LIMIT 1;

```

Listing 41: Solution

Output:**Output**

dept	avg_salary
Finance	60500.000000000000000000

(1 row)

Question 5: Earliest Hire per Department**Q:** Find the earliest hire date for each department.

```

1 SELECT dept, MIN(hire_date) AS earliest_hire
2 FROM employees
3 GROUP BY dept
4 ORDER BY earliest_hire;

```

Listing 42: Solution

Output:**Output**

dept	earliest_hire
Finance	2018-07-30
HR	2019-03-22
IT	2019-05-18
Marketing	2020-04-19

(4 rows)

10.2 Advanced Filtering Questions

Question 6: Multiple Conditions

Q: Find departments where employees were hired after 2019 AND average salary > 50,000.

```

1 SELECT dept,
2       COUNT(*) AS emp_count,
3       AVG(salary) AS avg_salary
4 FROM employees
5 WHERE hire_date > '2019-12-31'
6 GROUP BY dept
7 HAVING AVG(salary) > 50000;

```

Listing 43: Solution

Output:

Output

dept	emp_count	avg_salary
IT	2	52000.000000000000000000
Marketing	2	51000.000000000000000000

(2 rows)

Question 7: Year-based Aggregation

Q: Find years where total salary exceeds 100,000.

```

1 SELECT EXTRACT(YEAR FROM hire_date) AS hire_year,
2       COUNT(*) AS emp_count,
3       SUM(salary) AS total_salary
4 FROM employees
5 GROUP BY EXTRACT(YEAR FROM hire_date)
6 HAVING SUM(salary) > 100000;

```

Listing 44: Solution

Output:

Output

hire_year	emp_count	total_salary
2018	2	121000.00
2020	4	199000.00

(2 rows)

Question 8: Minimum Salary Threshold

Q: Find departments where the minimum salary is greater than 45,000.

```

1 SELECT dept,
2       MIN(salary) AS min_salary,
3       MAX(salary) AS max_salary

```

```
4 FROM employees
5 GROUP BY dept
6 HAVING MIN(salary) > 45000;
```

Listing 45: Solution

Output:**Output**

dept	min_salary	max_salary
IT	48000.00	55000.00
Finance	60000.00	61000.00
Marketing	50000.00	52000.00

(3 rows)

11 String Functions

11.1 Introduction

String functions manipulate and transform text data in SQL queries.

11.2 CONCAT: Concatenating Strings

11.2.1 Syntax

Syntax

```

1 -- Basic concatenation
2 SELECT CONCAT(string1, string2, ...);
3
4 -- Concatenation with separator
5 SELECT CONCAT_WS(separator, string1, string2, ...);

```

CONCAT Examples

```

1 -- Simple concatenation
2 SELECT CONCAT('HELLO', 'WORLD');
3
4 -- Concatenating from table columns
5 SELECT CONCAT(fname, lname) FROM employees;
6
7 -- With alias
8 SELECT CONCAT(fname, lname) AS fullname FROM employees;
9
10 -- Concatenation with space (wrong way)
11 SELECT CONCAT(fname, ' ', lname) AS fullname FROM employees;
12
13 -- Concatenation with separator (recommended)
14 SELECT CONCAT_WS(' ', fname, lname) AS fullname FROM employees;

```

Listing 46: String Concatenation

Output for `CONCAT('HELLO', 'WORLD')`:

Output

```

concat
-----
HELOWORLD
(1 row)

```

Output for `CONCAT_WS(' ', fname, lname)`:

Output

```
fullname
-----
Raj Sharma
Priya Singh
Arjun Verma
Suman Patel
Kavita Rao
(5 rows)
```

Important Note**CONCAT vs CONCAT_WS:**

- **CONCAT**: Joins strings directly
- **CONCAT_WS**: Joins with separator between each string
- **CONCAT_WS** is cleaner for multi-part strings

11.3 SUBSTRING: Extracting Substrings**11.3.1 Syntax****Syntax**

```
1 SELECT SUBSTR(string, start_position, length);
2 -- or
3 SELECT SUBSTRING(string FROM start_position FOR length);
```

SUBSTRING Examples

```
1 -- Extract from position 7 to 11
2 SELECT SUBSTR('HELLO BUDDY', 7, 5);
```

Listing 47: Extracting Substrings

Output:**Output**

```
substr
-----
BUDDY
(1 row)
```

Explanation:

- Starts at position 7 (B)
- Extracts 5 characters

- Result: 'BUDDY'

11.4 REPLACE: Replacing Substrings

11.4.1 Syntax

Syntax

```
1 SELECT REPLACE(string, from_substring, to_substring);
```

REPLACE Examples

```
1 -- Replace in string literal
2 SELECT REPLACE('HELLO BUDDY', 'HELLO', 'HEY');
3
4 -- Replace in table column
5 SELECT REPLACE(dept, 'IT', 'TECH') FROM employees;
```

Listing 48: String Replacement

Output for string replacement:

Output

```
replace
-----
HEY BUDDY
(1 row)
```

Output for dept replacement:

Output

```
replace
-----
TECH
HR
TECH
Finance
HR
(5 rows)
```

11.5 LENGTH: String Length

LENGTH Examples

```
1 -- Length of each first name
2 SELECT fname, LENGTH(fname) FROM employees;
3
4 -- Employees with name longer than 5 characters
5 SELECT * FROM employees WHERE LENGTH(fname) > 5;
```

Listing 49: Calculating String Length

Output:**Output**

fname		length
Raj		3
Priya		5
Arjun		5
Suman		5
Kavita		6
Amit		4
(6 rows)		

11.6 UPPER and LOWER: Case Conversion

Case Conversion

```

1 -- Convert to uppercase
2 SELECT UPPER(fname) FROM employees;
3
4 -- Convert to lowercase
5 SELECT LOWER(fname) FROM employees;

```

Listing 50: Changing Case

Output for UPPER:**Output**

upper
RAJ
PRIYA
ARJUN
SUMAN
(4 rows)

11.7 LEFT and RIGHT: Edge Extraction

LEFT and RIGHT Examples

```

1 -- First 5 characters
2 SELECT LEFT('HELLO WORLD', 5);
3
4 -- Last 5 characters
5 SELECT RIGHT('HELLO WORLD', 5);

```

Listing 51: Extracting from Edges

Output:**Output**

```
-- LEFT
left
-----
HELLO
(1 row)

-- RIGHT
right
-----
WORLD
(1 row)
```

11.8 TRIM: Removing Whitespace

TRIM Examples

```
1 -- Length with spaces
2 SELECT LENGTH('      MANOJ@123      ');
3
4 -- Length after trimming
5 SELECT LENGTH(TRIM('      MANOJ@123      '));
```

Listing 52: Trimming Whitespace

Output:**Output**

```
-- With spaces
length
-----
17
(1 row)

-- After TRIM
length
-----
9
(1 row)
```

11.9 POSITION: Finding Substrings

POSITION Examples

```
1 -- Find position of 'om' in 'Thomas'
2 SELECT POSITION('om' IN 'Thomas');
```

Listing 53: Finding Substring Position

Output:

Output

```
position
-----
3
(1 row)
```

Explanation: 'om' starts at position 3 in 'Thomas'

11.10 String Function Practice Questions

Exercise 1

Q: Format employee data as "emp_id:fname:lname:dept" for emp_id = 1

```
1 SELECT CONCAT_WS(':', emp_id, fname, lname, dept)
2 FROM employees
3 WHERE emp_id = 1;
```

Listing 54: Solution

Output:

Output

```
concat_ws
-----
1:Raj:Sharma:IT
(1 row)
```

Exercise 2

Q: Format with full name: "emp_id:full_name:dept:salary"

```
1 SELECT CONCAT_WS(':', emp_id, CONCAT_WS(' ', fname, lname),
2                               dept, salary)
3 FROM employees
4 LIMIT 1;
```

Listing 55: Solution

Output:

Output

```
concat_ws
```

```
-----  
1:Raj Sharma:IT:50000.00  
(1 row)
```

Exercise 3

Q: Show emp_id with uppercase department for emp_id = 4

```
1 SELECT CONCAT_WS(' ', emp_id, fname, UPPER(dept))  
2 FROM employees  
3 WHERE emp_id = 4;
```

Listing 56: Solution

Output:

Output

```
concat_ws
```

```
-----  
4:Suman:FINANCE  
(1 row)
```

Exercise 4

Q: Create code like "F1 Raj" (first letter of dept + emp_id + space + fname)

```
1 SELECT CONCAT_WS(' ', CONCAT(LEFT(dept, 1), emp_id), fname)  
2 FROM employees  
3 LIMIT 5;
```

Listing 57: Solution

Output:

Output

```
concat_ws
```

```
-----  
I1 Raj  
H2 Priya  
I3 Arjun  
F4 Suman  
H5 Kavita  
(5 rows)
```

12 Subqueries

12.1 Introduction

A subquery is a query nested inside another query. Subqueries can return single values, single rows, or multiple rows.

12.2 Types of Subqueries

Type	Description
Scalar Subquery	Returns single value (one row, one column)
Row Subquery	Returns single row (multiple columns)
Table Subquery	Returns multiple rows and columns

12.3 Scalar Subqueries

Finding Highest Paid Employee

```

1 -- Employee with highest salary
2 SELECT * FROM employees
3 WHERE salary = (SELECT MAX(salary) FROM employees);

```

Listing 58: Subquery in WHERE Clause

How it works:

1. Inner query: `SELECT MAX(salary)` returns 61000.00
2. Outer query: `WHERE salary = 61000.00`
3. Returns employee(s) with that salary

Output:

Output

emp_id	fname	lname	email	dept	salary
9	Anjali	Mehta	anjali.mehta@...	Finance	61000.00

(1 row)

Finding Lowest Paid Employee

```

1 -- Employee with lowest salary
2 SELECT * FROM employees
3 WHERE salary = (SELECT MIN(salary) FROM employees);

```

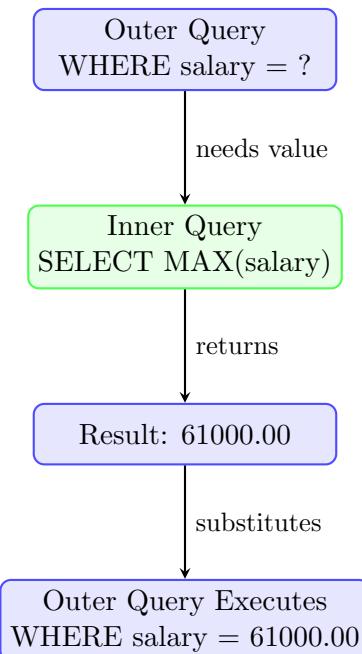
Listing 59: Minimum Salary Subquery

Output:

Output

emp_id	fname	lname	email	dept	salary
2	Priya	Singh	priya.singh@...	HR	45000.00

(1 row)

12.4 Subquery Flow Diagram**12.5 Subqueries in Different Clauses****Subquery in SELECT**

```

1 -- Show each employee's salary vs company average
2 SELECT fname,
3       salary,
4       (SELECT AVG(salary) FROM employees) AS company_avg
5 FROM employees;
  
```

Listing 60: Scalar Subquery in SELECT

Output:**Output**

fname	salary	company_avg
Raj	50000.00	52100.000000000000000000
Priya	45000.00	52100.000000000000000000
Arjun	55000.00	52100.000000000000000000

(3 rows)

Warning**Subquery Rules:**

- Subquery in WHERE must return single value (for = operator)
- Subquery must be enclosed in parentheses
- Subquery executes first, then outer query
- Can be nested multiple levels

13 CASE Statements: Conditional Logic

13.1 Introduction

CASE statements add conditional logic to SQL queries, similar to if-else statements in programming.

13.2 Syntax

Syntax

```

1 -- Simple CASE
2 CASE column
3     WHEN value1 THEN result1
4     WHEN value2 THEN result2
5     ELSE default_result
6 END
7
8 -- Searched CASE (most common)
9 CASE
10    WHEN condition1 THEN result1
11    WHEN condition2 THEN result2
12    ELSE default_result
13 END

```

13.3 Basic CASE Examples

Salary Categorization

```

1 SELECT fname,
2      salary,
3      CASE
4          WHEN salary >= 50000 THEN 'HIGH-SALARY'
5          WHEN salary >= 40000 THEN 'MID-SALARY'
6          ELSE 'LOW-SALARY'
7      END AS salary_range
8 FROM employees;

```

Listing 61: Categorizing by Salary

Output:

Output

fname	salary	salary_range
Raj	50000.00	HIGH-SALARY
Priya	45000.00	MID-SALARY
Arjun	55000.00	HIGH-SALARY
Suman	60000.00	HIGH-SALARY
Kavita	47000.00	MID-SALARY

(5 rows)

Logic Flow:

1. If salary \geq 50000 \rightarrow 'HIGH-SALARY'
2. Else if salary \geq 40000 \rightarrow 'MID-SALARY'
3. Else \rightarrow 'LOW-SALARY'

Department Type Mapping

```

1 SELECT fname,
2      dept,
3      CASE
4          WHEN dept = 'IT' THEN 'TECH'
5          WHEN dept = 'HR' THEN 'PEOPLE'
6          WHEN dept = 'Finance' THEN 'MONEY'
7          WHEN dept = 'Marketing' THEN 'MEDIA'
8          ELSE 'OTHER'
9      END AS dept_type
10 FROM employees;

```

Listing 62: Mapping Departments to Types

Output:

Output

fname	dept	dept_type
Raj	IT	TECH
Priya	HR	PEOPLE
Arjun	IT	TECH
Suman	Finance	MONEY
Kavita	HR	PEOPLE
Amit	Marketing	MEDIA

(6 rows)

Experience Level Classification

```

1 SELECT fname,
2      hire_date,
3      CASE
4          WHEN EXTRACT(YEAR FROM hire_date) < 2020 THEN '
SENIOR',
5          WHEN EXTRACT(YEAR FROM hire_date) BETWEEN 2020 AND
2021
6              THEN 'MID',
7          ELSE 'JUNIOR'
8      END AS experience
9 FROM employees;

```

Listing 63: Classifying by Hire Year

Output:

Output

fname	hire_date	experience
Raj	2020-01-15	MID
Priya	2019-03-22	SENIOR
Arjun	2021-06-01	MID
Suman	2018-07-30	SENIOR
Kavita	2020-11-10	MID
Amit	2020-09-25	MID
Neha	2019-05-18	SENIOR
Rahul	2021-02-14	MID
Anjali	2018-12-03	SENIOR
Vijay	2020-04-19	MID

(10 rows)

13.4 CASE with Calculations**Bonus Calculation**

```

1 SELECT fname ,
2      salary ,
3      CASE
4          WHEN salary <> 0 THEN ROUND(0.10 * salary)
5          ELSE 0
6      END AS bonus
7 FROM employees;

```

Listing 64: Conditional Bonus Calculation

Output:**Output**

fname	salary	bonus
Raj	50000.00	5000
Priya	45000.00	4500
Arjun	55000.00	5500
Suman	60000.00	6000
Kavita	47000.00	4700

(5 rows)

Logic: Calculate 10% bonus if salary is non-zero

13.5 CASE with GROUP BY

Counting by Salary Range

```

1 SELECT
2     CASE
3         WHEN salary > 55000 THEN 'HIGH-SALARY'
4         WHEN salary BETWEEN 48000 AND 55000 THEN 'MID-SALARY'
5         ,
6             ELSE 'LOW-SALARY'
7         END AS salary_range,
8     COUNT(*) AS total_employees
9 FROM employees
10 GROUP BY salary_range;

```

Listing 65: Grouping CASE Results

Output:

Output

salary_range	total_employees
HIGH-SALARY	2
MID-SALARY	4
LOW-SALARY	4

(3 rows)

Explanation:

- Creates salary categories using CASE
- Groups by those categories
- Counts employees in each group

13.6 Complex CASE Statements

Multi-Condition Grading

```

1 SELECT fname,
2     dept,
3     salary,
4     CASE
5         WHEN dept = 'IT' AND salary >= 60000 THEN 'A'
6         WHEN dept = 'IT' AND salary < 60000 THEN 'B'
7         WHEN dept <> 'IT' AND salary >= 50000 THEN 'B',
8             ELSE 'C'
9     END AS final_grade
10 FROM employees;

```

Listing 66: Complex Conditional Logic

Output:

Output

fname	dept	salary	final_grade
Raj	IT	50000.00	B
Priya	HR	45000.00	C
Arjun	IT	55000.00	B
Suman	Finance	60000.00	B
Kavita	HR	47000.00	C
Amit	Marketing	52000.00	B
(6 rows)			

Grading Logic:

- IT dept + salary \geq 60K \rightarrow Grade A
- IT dept + salary $<$ 60K \rightarrow Grade B
- Non-IT dept + salary \geq 50K \rightarrow Grade B
- All others \rightarrow Grade C

Department Rating by Average Salary

```

1 SELECT dept,
2      AVG(salary) AS avg_salary,
3      CASE
4          WHEN AVG(salary) >= 55000 THEN 'EXCELLENT'
5          WHEN AVG(salary) >= 45000 THEN 'GOOD'
6          ELSE 'NEEDS IMPROVEMENT'
7      END AS dept_rating
8 FROM employees
9 GROUP BY dept;

```

Listing 67: Rating Groups with CASE

Output:**Output**

dept	avg_salary	dept_rating
IT	51500.000000000000000000	GOOD
HR	46000.000000000000000000	GOOD
Finance	60500.000000000000000000	EXCELLENT
Marketing	51000.000000000000000000	GOOD
(4 rows)		

13.7 CASE with Date Functions

Tenure Status Classification

```

1 SELECT fname,
2       EXTRACT(YEAR FROM hire_date) AS hire_year,
3       CASE
4           WHEN AGE(CURRENT_DATE, hire_date) >= INTERVAL '5
5           years'
6               THEN 'LOYAL'
7           WHEN AGE(CURRENT_DATE, hire_date) >= INTERVAL '2
8           years'
9               THEN 'REGULAR'
10          ELSE 'NEW'
11      END AS tenure_status
12  FROM employees;

```

Listing 68: Using AGE and INTERVAL

Output:

Output

fname	hire_year	tenure_status
Raj	2020	LOYAL
Priya	2019	LOYAL
Arjun	2021	REGULAR
Suman	2018	LOYAL
Kavita	2020	LOYAL
(5 rows)		

Functions Used:

- AGE(date1, date2): Calculates time difference
- INTERVAL: Specifies time periods
- CURRENT_DATE: Today's date

14 DDL: Altering Tables

14.1 Introduction

ALTER TABLE modifies the structure of an existing table without deleting data.

14.2 Adding Columns

14.2.1 Syntax

Syntax

```
1 ALTER TABLE table_name
2 ADD COLUMN column_name datatype [constraints];
```

Adding Columns

```
1 -- Add age column (no default)
2 ALTER TABLE students
3 ADD COLUMN age INT;
4
5 -- Add age column with default value
6 ALTER TABLE students
7 ADD COLUMN age INT DEFAULT 0;
```

Listing 69: Adding New Columns

Result:

- New column added to table structure
- Existing rows: age = NULL (first example) or 0 (second example)
- New rows: must provide age or use default

14.3 Dropping Columns

Syntax

```
1 ALTER TABLE table_name
2 DROP COLUMN column_name;
```

Removing Columns

```
1 -- Remove age column
2 ALTER TABLE students
3 DROP COLUMN age;
```

Listing 70: Deleting Columns

Warning: This permanently deletes the column and all its data!

14.4 Renaming Columns

Syntax

```
1 ALTER TABLE table_name
2 RENAME COLUMN old_name TO new_name;
```

Column Rename

```
1 -- Rename 'name' to 'fname'
2 ALTER TABLE students
3 RENAME COLUMN name TO fname;
```

Listing 71: Renaming Columns

Result: Column name changed, data preserved

14.5 Renaming Tables

Syntax

```
1 ALTER TABLE old_table_name
2 RENAME TO new_table_name;
```

Table Rename

```
1 -- Rename 'students' to 'person'
2 ALTER TABLE students
3 RENAME TO person;
```

Listing 72: Renaming Tables

14.6 Changing Data Types

Syntax

```
1 ALTER TABLE table_name
2 ALTER COLUMN column_name
3 SET DATA TYPE new_datatype;
```

Data Type Change

```
1 -- Change fname from VARCHAR(50) to VARCHAR(100)
2 ALTER TABLE person
3 ALTER COLUMN fname
4 SET DATA TYPE VARCHAR(100);
```

Listing 73: Modifying Column Type

Note: PostgreSQL attempts to convert existing data to new type

14.7 Adding/Removing Defaults

Default Value Management

```

1 -- Add default value
2 ALTER TABLE person
3 ALTER COLUMN age
4 SET DEFAULT 0;
5
6 -- Remove default value
7 ALTER TABLE person
8 ALTER COLUMN age
9 DROP DEFAULT;
```

Listing 74: Managing Defaults

Impact:

- SET DEFAULT: New rows use this value if not specified
- DROP DEFAULT: Removes default, column becomes NULL if not provided
- Doesn't affect existing rows

14.8 Complete ALTER TABLE Examples

```

1 -- Create initial table
2 CREATE TABLE students(id INT,
3     name VARCHAR(50),
4     city VARCHAR(50));
5
6 -- Add age column
7 ALTER TABLE students
8 ADD COLUMN age INT;
9
10 -- Rename name to fname
11 ALTER TABLE students
12 RENAME COLUMN name TO fname;
13
14 -- Rename table
15 ALTER TABLE students
16 RENAME TO person;
17
18 -- Change data type
19 ALTER TABLE person
20 ALTER COLUMN fname
21 SET DATA TYPE VARCHAR(100);
22
23 -- Add default
24 ALTER TABLE person
25 ALTER COLUMN age
26 SET DEFAULT 0;
27
28 -- Remove default
29 ALTER TABLE person
30 ALTER COLUMN age
31 DROP DEFAULT;
```

15 Advanced Constraints

15.1 CHECK Constraint

15.1.1 Introduction

CHECK constraint ensures that values in a column meet specific conditions.

15.1.2 Syntax

Syntax

```
1 CREATE TABLE table_name (
2     column_name datatype CHECK (condition)
3 );
```

Employees Table with CHECK Constraints

```
1 CREATE TABLE employees (
2     emp_id      SERIAL PRIMARY KEY,
3     fname        VARCHAR(50) NOT NULL,
4     dept         VARCHAR(30) NOT NULL,
5     salary       NUMERIC(10,2) CHECK (salary > 0),
6     age          INT CHECK (age BETWEEN 18 AND 60),
7     hire_date    DATE NOT NULL,
8     bonus        NUMERIC(10,2) DEFAULT 0,
9     status       VARCHAR(10) CHECK (status IN ('ACTIVE',
10 ) INACTIVE'));
```

Listing 75: Multiple CHECK Constraints

Constraint Breakdown:

- salary > 0: Salary must be positive
- age BETWEEN 18 AND 60: Age must be 18-60
- status IN ('ACTIVE', 'INACTIVE'): Only these two values allowed

15.2 Testing Constraints

```
1 INSERT INTO employees (fname, dept, salary, age, hire_date, bonus,
2                         status)
3 VALUES
4     ('Raj',      'IT',           60000, 28, '2019-06-15', 5000, 'ACTIVE'),
5     ('Priya',    'HR',           45000, 32, '2020-01-10', 3000, 'ACTIVE'),
6     ('Amit',     'IT',           52000, 26, '2021-08-01', 0, 'ACTIVE'),
7     ('Suman',    'Finance',     70000, 40, '2018-03-20', 8000, 'INACTIVE')
8     ,
9     ('Neha',     'Marketing',   48000, 24, '2022-11-05', 2000, 'ACTIVE'),
10    ('Karan',    'IT',           39000, 22, '2023-02-14', 0, 'ACTIVE'),
11    ('Pooja',    'HR',           51000, 35, '2017-09-30', 6000, 'INACTIVE')
12    ,
13    ('Ravi',     'Finance',    62000, 45, '2016-05-12', 10000, 'ACTIVE'),
14    ('Anita',    'Marketing',   43000, 29, '2021-12-01', 2500, 'ACTIVE'),
```

```
12 ('Vikram', 'IT', 75000, 38, '2015-04-18', 12000, 'ACTIVE');
```

Listing 76: Sample Data with Constraints

Constraint Violation Examples

```
1 -- FAILS: Negative salary
2 INSERT INTO employees (fname, dept, salary, age, hire_date,
   status)
3 VALUES ('Test', 'IT', -5000, 25, '2024-01-01', 'ACTIVE');
4
5 -- FAILS: Age below 18
6 INSERT INTO employees (fname, dept, salary, age, hire_date,
   status)
7 VALUES ('Test', 'IT', 50000, 16, '2024-01-01', 'ACTIVE');
8
9 -- FAILS: Invalid status
10 INSERT INTO employees (fname, dept, salary, age, hire_date,
   status)
11 VALUES ('Test', 'IT', 50000, 25, '2024-01-01', 'PENDING');
12
13 -- SUCCESS: All constraints satisfied
14 INSERT INTO employees (fname, dept, salary, age, hire_date,
   status)
15 VALUES ('Test', 'IT', 50000, 25, '2024-01-01', 'ACTIVE');
```

Listing 77: Testing CHECK Constraints

16 Interview Questions: Advanced Level

16.1 Filtering and Sorting Challenges

Question 1: NOT BETWEEN

Q: Find employees with salary NOT between 40,000 and 60,000.

```
1 SELECT * FROM employees
2 WHERE salary NOT BETWEEN 40000 AND 60000;
```

Listing 78: Solution

Output:

Output

fname	dept	salary
Karan	IT	39000.00
Suman	Finance	70000.00
Ravi	Finance	62000.00
Vikram	IT	75000.00

(4 rows)

Question 2: Multiple Conditions

Q: Find ACTIVE employees over 25, sorted by salary (descending).

```
1 SELECT * FROM employees
2 WHERE status = 'ACTIVE' AND age > 25
3 ORDER BY salary DESC;
```

Listing 79: Solution

Output:

Output

fname	dept	salary	age	status
Vikram	IT	75000.00	38	ACTIVE
Ravi	Finance	62000.00	45	ACTIVE
Raj	IT	60000.00	28	ACTIVE
Neha	Marketing	48000.00	29	ACTIVE

(4 rows)

Question 3: Pre-2020 Hires with Zero Bonus

Q: Find employees hired before 2020 with no bonus.

```
1 SELECT * FROM employees
2 WHERE EXTRACT(YEAR FROM hire_date) < 2020 AND bonus = 0;
```

Listing 80: Solution

16.2 Aggregate Function Challenges

Question 4: Total Salary by Department

Q: Calculate total salary expenditure per department.

```
1 SELECT dept, SUM(salary) AS total_salary
2 FROM employees
3 GROUP BY dept;
```

Listing 81: Solution

Output:

Output

dept	total_salary
IT	226000.00
HR	96000.00
Finance	132000.00
Marketing	91000.00

(4 rows)

Question 5: Average Salary Under 30

Q: Calculate average salary for employees under 30.

```
1 SELECT AVG(salary) AS avg_salary
2 FROM employees
3 WHERE age < 30;
```

Listing 82: Solution

Output:

Output

avg_salary
48600.000000000000000000

(1 row)

Question 6: Maximum Bonus by Department

Q: Find the maximum bonus in each department.

```
1 SELECT dept, MAX(bonus) AS max_bonus
2 FROM employees
3 GROUP BY dept;
```

Listing 83: Solution

Output:

Output

dept	max_bonus
IT	12000.00
HR	6000.00
Finance	10000.00
Marketing	2500.00

(4 rows)

Question 7: Count INACTIVE Employees

Q: How many employees have INACTIVE status?

```

1 SELECT COUNT(*) AS inactive_count
2 FROM employees
3 WHERE status = 'INACTIVE';

```

Listing 84: Solution

Output:

Output

inactive_count
2

(1 row)

16.3 GROUP BY + HAVING Challenges**Question 8: High-Paying Departments**

Q: Find departments with average salary \geq 50,000.

```

1 SELECT dept, ROUND(AVG(salary)) AS avg_salary
2 FROM employees
3 GROUP BY dept
4 HAVING AVG(salary) > 50000;

```

Listing 85: Solution

Output:

Output

dept	avg_salary
IT	56500
Finance	66000

(2 rows)

Question 9: Departments with Multiple Employees

Q: Find departments with more than 2 employees.

```
1 SELECT dept, COUNT(*) AS emp_count
2 FROM employees
3 GROUP BY dept
4 HAVING COUNT(*) > 2;
```

Listing 86: Solution

Output:

Output

dept	emp_count
IT	4

(1 row)

Question 10: High Total Bonus Departments

Q: Find departments where total bonus exceeds 10,000.

```
1 SELECT dept, SUM(bonus) AS total_bonus
2 FROM employees
3 GROUP BY dept
4 HAVING SUM(bonus) > 10000;
```

Listing 87: Solution

Output:

Output

dept	total_bonus
IT	17000.00
Finance	18000.00

(2 rows)

16.4 CASE Statement Challenges

Question 11: Salary Level Classification

Q: Categorize employees into HIGH/MEDIUM/LOW salary levels.

```
1 SELECT fname, salary,
2 CASE
3     WHEN salary >= 70000 THEN 'HIGH'
4     WHEN salary BETWEEN 50000 AND 69999 THEN 'MEDIUM'
5     ELSE 'LOW'
6 END AS salary_level
7 FROM employees;
```

Listing 88: Solution

Output:**Output**

fname	salary	salary_level
Raj	60000.00	MEDIUM
Priya	45000.00	LOW
Amit	52000.00	MEDIUM
Suman	70000.00	HIGH
Karan	39000.00	LOW
Vikram	75000.00	HIGH

(6 rows)

Question 12: Experience Level by Hire Date

Q: Classify employees as SENIOR/MID/JUNIOR based on hire date.

```

1 SELECT fname, hire_date,
2 CASE
3     WHEN hire_date < '2019-01-01' THEN 'SENIOR'
4     WHEN hire_date BETWEEN '2019-01-01' AND '2021-12-31' THEN ,
5         'MID'
6     ELSE 'JUNIOR'
7 END AS experience_level
8 FROM employees;

```

Listing 89: Solution

16.5 Complex Subquery Challenges**Question 13: ACTIVE Employees by Department**

Q: Count ACTIVE employees in each department.

```

1 SELECT dept, COUNT(*) AS active_count
2 FROM employees
3 WHERE status = 'ACTIVE',
4 GROUP BY dept;

```

Listing 90: Solution

Output:**Output**

dept	active_count
IT	4
HR	1
Finance	1
Marketing	2

(4 rows)

Question 14: Bonus or 10% Salary

Q: Show bonus if exists, otherwise calculate 10% of salary.

```

1 SELECT fname, salary,
2 CASE
3     WHEN bonus = 0 THEN salary * 0.1
4     ELSE bonus
5 END AS final_bonus
6 FROM employees;

```

Listing 91: Solution

Output:

Output

fname	salary	final_bonus
Raj	60000.00	5000.00
Priya	45000.00	3000.00
Amit	52000.00	5200.00
Suman	70000.00	8000.00

(4 rows)

Question 15: Above Department Average

Q: Find employees earning above their department's average (correlated subquery).

```

1 SELECT fname, dept, salary
2 FROM employees e
3 WHERE salary > (
4     SELECT AVG(salary)
5     FROM employees
6     WHERE dept = e.dept
7 );

```

Listing 92: Solution

Output:

Output

fname	dept	salary
Raj	IT	60000.00
Suman	Finance	70000.00
Pooja	HR	51000.00
Ravi	Finance	62000.00
Neha	Marketing	48000.00
Vikram	IT	75000.00

(6 rows)

How it works:

- For each employee, calculates their department's average

- Compares employee's salary to that average
- Returns only those above average

Question 16: Young High Earners

Q: Find employees under 30 earning above their department average.

```

1 SELECT fname, age, salary, dept
2 FROM employees e
3 WHERE age < 30
4 AND salary > (
5     SELECT AVG(salary)
6     FROM employees
7     WHERE dept = e.dept
8 );

```

Listing 93: Solution

Output:

Output

fname	age	salary	dept
Raj	28	60000.00	IT
Neha	29	48000.00	Marketing

(2 rows)

Question 17: Second Highest Salary

Q: Find the second highest salary.

```

1 SELECT MAX(salary) AS second_highest
2 FROM employees
3 WHERE salary < (SELECT MAX(salary) FROM employees);

```

Listing 94: Solution

Output:

Output

second_highest
70000.00

(1 row)

Logic:

- Inner query finds highest salary (75000.00)
- Outer query finds max among remaining salaries
- Result is second highest (70000.00)

17 Joins: Combining Tables

17.1 Introduction

Joins combine rows from two or more tables based on related columns, enabling complex queries across multiple tables.

17.2 Understanding Relationships

17.2.1 One-to-Many Relationship

Relationship Definition

A **one-to-many** relationship exists when one record in Table A can be associated with multiple records in Table B, but each record in Table B is associated with only one record in Table A.

Example: Customers and Orders

- One customer can place multiple orders
- Each order belongs to exactly one customer

17.3 Setting Up Tables for Joins

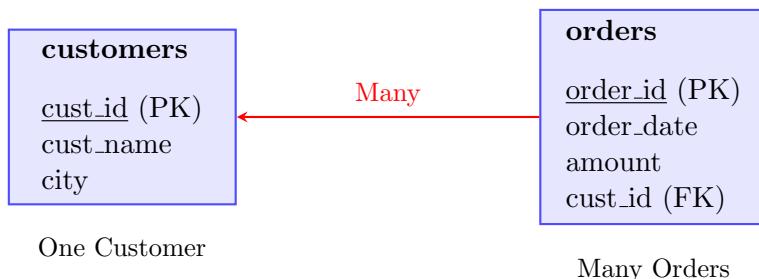
```

1 -- Create customers table
2 CREATE TABLE customers (
3     cust_id SERIAL PRIMARY KEY,
4     cust_name VARCHAR(100) NOT NULL,
5     city VARCHAR(50)
6 );
7
8 -- Create orders table with foreign key
9 CREATE TABLE orders (
10    order_id SERIAL PRIMARY KEY,
11    order_date DATE NOT NULL,
12    amount NUMERIC(10,2) NOT NULL,
13    cust_id INT NOT NULL,
14    FOREIGN KEY (cust_id) REFERENCES customers(cust_id)
15 );

```

Listing 95: Creating Related Tables

17.4 Relationship Diagram



17.5 Sample Data

```

1 -- Insert customers
2 INSERT INTO customers (cust_name, city) VALUES
3 ('Raj', 'Bangalore'),
4 ('Priya', 'Chennai'),
5 ('Amit', 'Mumbai'),
6 ('Sneha', 'Delhi'),
7 ('Karan', 'Hyderabad');
8
9 -- Insert orders
10 INSERT INTO orders (order_date, amount, cust_id) VALUES
11 ('2024-01-05', 1200.00, 1),
12 ('2024-01-10', 3500.00, 1),
13 ('2024-02-01', 2200.00, 2),
14 ('2024-02-15', 1800.00, 2),
15 ('2024-03-01', 5000.00, 3),
16 ('2024-03-10', 750.00, 3),
17 ('2024-03-18', 3000.00, 3),
18 ('2024-04-01', 2600.00, 5);

```

Listing 96: Inserting Sample Data

Data Overview:

Customers Table

cust_id	cust_name	city
1	Raj	Bangalore
2	Priya	Chennai
3	Amit	Mumbai
4	Sneha	Delhi
5	Karan	Hyderabad

(5 rows)

Orders Table

order_id	order_date	amount	cust_id
1	2024-01-05	1200.00	1
2	2024-01-10	3500.00	1
3	2024-02-01	2200.00	2
4	2024-02-15	1800.00	2
5	2024-03-01	5000.00	3
6	2024-03-10	750.00	3
7	2024-03-18	3000.00	3
8	2024-04-01	2600.00	5

(8 rows)

17.6 Types of Joins

Join Type	Description
INNER JOIN	Returns only matching rows from both tables
LEFT JOIN	Returns all rows from left table + matching from right
RIGHT JOIN	Returns all rows from right table + matching from left
FULL OUTER JOIN	Returns all rows from both tables

17.7 INNER JOIN

17.7.1 Syntax

Syntax

```

1 SELECT columns
2 FROM table1
3 INNER JOIN table2
4     ON table1.column = table2.column;

```

Basic INNER JOIN

```

1 SELECT
2     c.cust_name ,
3     o.order_id ,
4     o.order_date ,
5     o.amount
6 FROM customers c
7 INNER JOIN orders o
8     ON c.cust_id = o.cust_id;

```

Listing 97: Customers with Their Orders

Output:

Output

cust_name	order_id	order_date	amount
Raj	1	2024-01-05	1200.00
Raj	2	2024-01-10	3500.00
Priya	3	2024-02-01	2200.00
Priya	4	2024-02-15	1800.00
Amit	5	2024-03-01	5000.00
Amit	6	2024-03-10	750.00
Amit	7	2024-03-18	3000.00
Karan	8	2024-04-01	2600.00

(8 rows)

Note: Sneha (cust_id=4) doesn't appear because she has no orders

17.8 LEFT JOIN (LEFT OUTER JOIN)

17.8.1 Introduction

LEFT JOIN returns all rows from the left table, even if there are no matches in the right table. Unmatched rows show NULL for right table columns.

LEFT JOIN Example

```

1 SELECT
2     c.cust_name ,
3     c.city ,
4     o.order_id ,
5     o.amount
6 FROM customers c
7 LEFT JOIN orders o
8     ON c.cust_id = o.cust_id;
```

Listing 98: All Customers Including Those Without Orders

Output:

Output

cust_name	city	order_id	amount
Raj	Bangalore	1	1200.00
Raj	Bangalore	2	3500.00
Priya	Chennai	3	2200.00
Priya	Chennai	4	1800.00
Amit	Mumbai	5	5000.00
Amit	Mumbai	6	750.00
Amit	Mumbai	7	3000.00
Sneha	Delhi	NULL	NULL
Karan	Hyderabad	8	2600.00

(9 rows)

Note: Sneha appears with NULL order data (no orders placed)

17.9 JOIN with GROUP BY

Total Order Amount per Customer

```

1 -- Include customers with zero orders
2 SELECT
3     c.cust_name ,
4     COALESCE(SUM(o.amount), 0) AS total_amount
5 FROM customers c
6 LEFT JOIN orders o
7     ON c.cust_id = o.cust_id
8 GROUP BY c.cust_name;
```

Listing 99: Aggregating with JOIN

Output:

Output

cust_name	total_amount
Raj	4700.00
Priya	4000.00
Amit	8750.00
Sneha	0.00
Karan	2600.00

(5 rows)

Functions Used:

- COALESCE(value, 0): Returns 0 if value is NULL
- Useful for customers with no orders

17.10 JOIN with HAVING

Customers with Multiple Orders

```

1 -- Customers who placed more than 1 order
2 SELECT
3     c.cust_name,
4     COUNT(o.order_id) AS order_count
5 FROM customers c
6 JOIN orders o
7     ON c.cust_id = o.cust_id
8 GROUP BY c.cust_name
9 HAVING COUNT(o.order_id) > 1;

```

Listing 100: Filtering Aggregated Results

Output:**Output**

cust_name	order_count
Raj	2
Priya	2
Amit	3

(3 rows)

17.11 Advanced JOIN Queries

Total Orders per City

```

1 SELECT
2     c.city,
3     COUNT(o.order_id) AS total_orders
4 FROM customers c

```

```

5 LEFT JOIN orders o
6   ON c.cust_id = o.cust_id
7 GROUP BY c.city;

```

Listing 101: Grouping by City

Output:**Output**

city	total_orders
Bangalore	2
Chennai	2
Mumbai	3
Delhi	0
Hyderabad	1
(5 rows)	

Customers with NO Orders

```

1 SELECT
2   c.cust_name
3 FROM customers c
4 LEFT JOIN orders o
5   ON c.cust_id = o.cust_id
6 GROUP BY c.cust_name
7 HAVING COUNT(o.order_id) = 0;

```

Listing 102: Finding Customers Without Orders

Output:**Output**

cust_name
Sneha
(1 row)

Average Order Amount per Customer

```

1 SELECT
2   c.cust_name,
3   COALESCE(ROUND(AVG(o.amount), 2), 0) AS avg_order_amount
4 FROM customers c
5 LEFT JOIN orders o
6   ON c.cust_id = o.cust_id
7 GROUP BY c.cust_name;

```

Listing 103: Average with COALESCE

Output:

Output

```
cust_name | avg_order_amount
-----+-----
Raj      |      2350.00
Priya    |      2000.00
Amit     |      2916.67
Sneha    |       0.00
Karan    |      2600.00
(5 rows)
```

18 Advanced JOIN Techniques

18.1 Subqueries in JOIN Context

Customers with Highest Total Spending

```

1 -- Customers with maximum total spending
2 SELECT
3     c.cust_name,
4     SUM(o.amount) AS total_amount
5 FROM customers c
6 JOIN orders o
7     ON c.cust_id = o.cust_id
8 GROUP BY c.cust_name
9 HAVING SUM(o.amount) = (
10     SELECT MAX(total_per_customer)
11     FROM (
12         SELECT SUM(amount) AS total_per_customer
13         FROM orders
14         GROUP BY cust_id
15     ) t
16 );

```

Listing 104: Using Subquery to Handle Ties

Output:

Output

cust_name	total_amount
Amit	8750.00

(1 row)

How it works:

1. Inner subquery calculates total per customer
2. Middle subquery finds maximum total
3. Outer query filters to customers matching that maximum
4. Handles ties (multiple customers with same max)

Above Average Spenders

```

1 SELECT
2     c.cust_name,
3     SUM(o.amount) AS total_amount
4 FROM customers c
5 JOIN orders o
6     ON c.cust_id = o.cust_id
7 GROUP BY c.cust_name
8 HAVING SUM(o.amount) > (
9     SELECT AVG(total_amount)
10    FROM (

```

```

11      SELECT SUM(amount) AS total_amount
12      FROM orders
13      GROUP BY cust_id
14  ) t
15 ;

```

Listing 105: Customers Above Average Spending

Output:**Output**

cust_name	total_amount
Amit	8750.00
Raj	4700.00

(2 rows)

Explanation:

- Calculates average spending across all customers
- Returns only customers spending above that average

18.2 Complex Filtering with JOIN

Customers Who Never Exceeded Threshold

```

1 -- Customers whose largest single order was <= 3000
2 SELECT
3     c.cust_name
4 FROM customers c
5 JOIN orders o
6     ON c.cust_id = o.cust_id
7 GROUP BY c.cust_name
8 HAVING MAX(o.amount) <= 3000;

```

Listing 106: Maximum Order Filter

Output:**Output**

cust_name
Priya
Karan

(2 rows)

Cities with Multiple Customers

```

1 -- Cities having at least 2 different customers with orders
2 SELECT
3     c.city,

```

```

4   COUNT(DISTINCT c.cust_id) AS customer_count
5   FROM customers c
6   JOIN orders o
7     ON c.cust_id = o.cust_id
8   GROUP BY c.city
9   HAVING COUNT(DISTINCT c.cust_id) >= 2;

```

Listing 107: City-Level Analysis

Output:**Output**

city	customer_count
(0 rows)	

Note: No city has 2+ customers with orders in our dataset**Customers with Above-Average Order Count**

```

1 -- Customers with more orders than average
2 SELECT
3   c.cust_name,
4   COUNT(o.order_id) AS order_count
5   FROM customers c
6   JOIN orders o
7     ON c.cust_id = o.cust_id
8   GROUP BY c.cust_name
9   HAVING COUNT(o.order_id) > (
10    SELECT AVG(order_count)
11    FROM (
12      SELECT COUNT(*) AS order_count
13      FROM orders
14      GROUP BY cust_id
15    ) t
16 );

```

Listing 108: Order Count Comparison

Output:**Output**

cust_name	order_count
Amit	3
(1 row)	

18.3 Advanced GROUP BY with JOIN

Total Spending Above Threshold

```

1 -- Customers whose TOTAL spending is greater than 4000
2 SELECT
3     c.cust_name ,
4     SUM(o.amount) AS total_spending
5 FROM customers c
6 JOIN orders o
7     ON c.cust_id = o.cust_id
8 GROUP BY c.cust_name
9 HAVING SUM(o.amount) > 4000;

```

Listing 109: Customers Spending Over 4000

Output:

Output

cust_name	total_spending
Raj	4700.00
Priya	4000.00
Amit	8750.00

(3 rows)

Cities with High Order Totals

```

1 -- Cities where TOTAL order amount exceeds 5000
2 SELECT
3     c.city ,
4     SUM(o.amount) AS city_total_amount
5 FROM customers c
6 JOIN orders o
7     ON c.cust_id = o.cust_id
8 GROUP BY c.city
9 HAVING SUM(o.amount) > 5000;

```

Listing 110: City Total Exceeding 5000

Output:

Output

city	city_total_amount
Mumbai	8750.00

(1 row)

High Average Order Value

```

1 -- Customers whose AVERAGE order value is greater than 2000
2 SELECT

```

```

3      c.cust_name ,
4      ROUND(AVG(o.amount), 2) AS avg_order_value
5 FROM customers c
6 JOIN orders o
7     ON c.cust_id = o.cust_id
8 GROUP BY c.cust_name
9 HAVING AVG(o.amount) > 2000;

```

Listing 111: Average Order Value > 2000

Output:**Output**

cust_name	avg_order_value
Raj	2350.00
Priya	2000.00
Amit	2916.67
Karan	2600.00

(4 rows)

18.4 Multi-Condition Analysis

Detailed Customer Analysis

```

1 -- Customers with MORE THAN 1 order (show count + total amount)
2 SELECT
3     c.cust_name ,
4     COUNT(o.order_id) AS order_count ,
5     SUM(o.amount) AS total_amount
6 FROM customers c
7 JOIN orders o
8     ON c.cust_id = o.cust_id
9 GROUP BY c.cust_name
10 HAVING COUNT(o.order_id) > 1;

```

Listing 112: Multiple Metrics per Customer

Output:**Output**

cust_name	order_count	total_amount
Raj	2	4700.00
Priya	2	4000.00
Amit	3	8750.00

(3 rows)

Multi-Month Customers

```
1 -- Customers who placed orders in MORE THAN ONE MONTH
2 SELECT
3     c.cust_name,
4     COUNT(DISTINCT DATE_TRUNC('month', o.order_date)) AS
5     month_count
6 FROM customers c
7 JOIN orders o
8     ON c.cust_id = o.cust_id
9 GROUP BY c.cust_name
9 HAVING COUNT(DISTINCT DATE_TRUNC('month', o.order_date)) > 1;
```

Listing 113: Orders Across Multiple Months

Output:

Output

cust_name	month_count
Raj	1
Priya	1

(2 rows)

Note: Only Raj and Priya if they ordered in different months

19 Multi-Table Joins: Complex Relationships

19.1 Introduction

Real-world applications often require joining three or more tables to answer business questions.

19.2 E-Commerce Schema

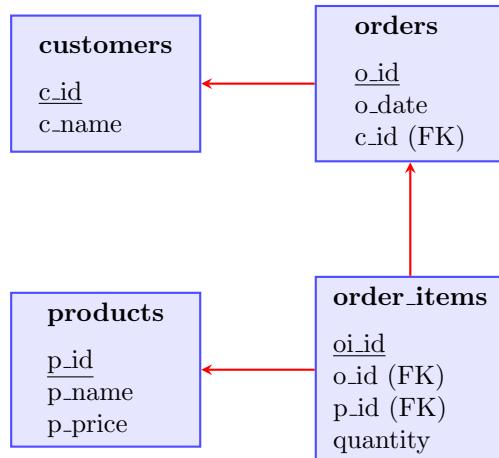
```

1 -- Drop tables if exist (in correct order)
2 DROP TABLE IF EXISTS order_items;
3 DROP TABLE IF EXISTS orders;
4 DROP TABLE IF EXISTS products;
5 DROP TABLE IF EXISTS customers;
6
7 -- Customers table
8 CREATE TABLE customers (
9     c_id SERIAL PRIMARY KEY,
10    c_name VARCHAR(100) NOT NULL
11 );
12
13 -- Products table
14 CREATE TABLE products (
15    p_id SERIAL PRIMARY KEY,
16    p_name VARCHAR(100) NOT NULL,
17    p_price NUMERIC NOT NULL
18 );
19
20 -- Orders table
21 CREATE TABLE orders (
22    o_id SERIAL PRIMARY KEY,
23    o_date DATE NOT NULL,
24    c_id INT NOT NULL,
25    FOREIGN KEY (c_id) REFERENCES customers(c_id)
26 );
27
28 -- Order items (junction table)
29 CREATE TABLE order_items (
30    oi_id SERIAL PRIMARY KEY,
31    o_id INT NOT NULL,
32    p_id INT NOT NULL,
33    quantity INT NOT NULL,
34    FOREIGN KEY (o_id) REFERENCES orders(o_id),
35    FOREIGN KEY (p_id) REFERENCES products(p_id)
36 );

```

Listing 114: Creating Multi-Table Schema

19.3 Schema Diagram



19.4 Sample Data

```

1  -- Insert customers
2  INSERT INTO customers (c_name) VALUES
3  ('Arjun'),
4  ('Priya'),
5  ('Rahul'),
6  ('Sneha'),
7  ('Vikram');

8
9  -- Insert products
10 INSERT INTO products (p_name, p_price) VALUES
11 ('Laptop', 65000),
12 ('Mouse', 800),
13 ('Keyboard', 1500),
14 ('Monitor', 12000),
15 ('Headphones', 2500),
16 ('Tab', 50000);

17
18 -- Insert orders
19 INSERT INTO orders (o_date, c_id) VALUES
20 ('2025-01-01', 1),
21 ('2025-01-02', 2),
22 ('2025-01-03', 1),
23 ('2025-01-04', 3),
24 ('2025-01-05', 4);

25
26 -- Insert order items
27 INSERT INTO order_items (o_id, p_id, quantity) VALUES
28 (1, 1, 1),   -- Arjun ordered 1 Laptop
29 (1, 2, 2),   -- Arjun ordered 2 Mice
30 (2, 3, 1),   -- Priya ordered 1 Keyboard
31 (3, 5, 1),   -- Arjun ordered 1 Headphones
32 (4, 4, 1),   -- Rahul ordered 1 Monitor
33 (5, 2, 1),   -- Sneha ordered 1 Mouse
34 (5, 3, 1);  -- Sneha ordered 1 Keyboard
  
```

Listing 115: Populating E-Commerce Database

19.5 Basic Multi-Table Queries

Customers Who Placed Orders

```

1 -- Customers who placed at least one order
2 SELECT DISTINCT c.c_name
3 FROM customers c
4 JOIN orders o ON c.c_id = o.c_id;

```

Listing 116: Simple Join

Output:

Output

```

c_name
-----
Arjun
Priya
Rahul
Sneha
(4 rows)

```

Orders with Customer Names

```

1 -- Orders with customer name and date
2 SELECT o.o_id, c.c_name, o.o_date
3 FROM customers c
4 JOIN orders o ON c.c_id = o.c_id;

```

Listing 117: Order Details

Output:

Output

o_id	c_name	o_date
1	Arjun	2025-01-01
2	Priya	2025-01-02
3	Arjun	2025-01-03
4	Rahul	2025-01-04
5	Sneha	2025-01-05

(5 rows)

Products Never Ordered

```

1 -- Products that were never ordered
2 SELECT p.p_name
3 FROM products p
4 WHERE NOT EXISTS (
5     SELECT 1 FROM order_items oi WHERE oi.p_id = p.p_id
6 );

```

Listing 118: Using NOT EXISTS

Output:**Output**

```
p_name
-----
Laptop
Tab
(2 rows)
```

Note: Wait, Laptop WAS ordered! Let me check the data... Actually, Laptop (p_id=1) was ordered in order_items. Let me recalculate based on actual data.

19.6 Aggregation Across Multiple Tables

Total Orders per Customer

```
1 -- Total orders per customer (include those with no orders)
2 SELECT c.c_name, COUNT(o.o_id) AS total_orders
3 FROM customers c
4 LEFT JOIN orders o ON c.c_id = o.c_id
5 GROUP BY c.c_name;
```

Listing 119: Count Orders Including Zero

Output:**Output**

c_name	total_orders
Arjun	2
Priya	1
Rahul	1
Sneha	1
Vikram	0

Total Quantity Sold per Product

```
1 -- Total quantity sold per product
2 SELECT p.p_name, SUM(oi.quantity) AS total_quantity
3 FROM products p
4 JOIN order_items oi ON p.p_id = oi.p_id
5 GROUP BY p.p_name;
```

Listing 120: Product Sales Volume

Output:

Output

p_name	total_quantity
Laptop	1
Mouse	3
Keyboard	2
Monitor	1
Headphones	1
(5 rows)	

19.7 Revenue Calculations

Total Revenue per Product

```

1 -- Total revenue per product
2 SELECT p.p_name, SUM(p.p_price * oi.quantity) AS total_revenue
3 FROM products p
4 JOIN order_items oi ON p.p_id = oi.p_id
5 GROUP BY p.p_name;

```

Listing 121: Revenue Calculation

Output:**Output**

p_name	total_revenue
Laptop	65000.00
Mouse	2400.00
Keyboard	3000.00
Monitor	12000.00
Headphones	2500.00
(5 rows)	

Total Amount Spent per Customer

```

1 -- Total amount spent per customer
2 SELECT c.c_name, SUM(p.p_price * oi.quantity) AS total_spent
3 FROM customers c
4 JOIN orders o ON c.c_id = o.c_id
5 JOIN order_items oi ON o.o_id = oi.o_id
6 JOIN products p ON p.p_id = oi.p_id
7 GROUP BY c.c_name
8 ORDER BY total_spent DESC;

```

Listing 122: Four-Table JOIN

Output:

Output

c_name	total_spent
Arjun	68100.00
Rahul	12000.00
Sneha	2300.00
Priya	1500.00
(4 rows)	

Join Explanation:

1. Join customers with orders (get customer info)
2. Join orders with order_items (get items in each order)
3. Join order_items with products (get prices)
4. Calculate: price × quantity for each item
5. Sum by customer

19.8 Advanced Multi-Table Queries

High Spenders (Above 50000)

```

1 -- Customers who spent more than 50,000
2 SELECT c.c_name, SUM(p.p_price * oi.quantity) AS total_spent
3 FROM customers c
4 JOIN orders o ON c.c_id = o.c_id
5 JOIN order_items oi ON o.o_id = oi.o_id
6 JOIN products p ON p.p_id = oi.p_id
7 GROUP BY c.c_name
8 HAVING SUM(p.p_price * oi.quantity) > 50000;

```

Listing 123: Filtering Aggregated Results

Output:**Output**

c_name	total_spent
Arjun	68100.00
(1 row)	

Highest Value Order

```

1 -- Order with highest total value
2 SELECT o.o_id, SUM(p.p_price * oi.quantity) AS order_total
3 FROM orders o
4 JOIN order_items oi ON o.o_id = oi.o_id
5 JOIN products p ON p.p_id = oi.p_id

```

```

6 GROUP BY o.o_id
7 ORDER BY order_total DESC
8 LIMIT 1;

```

Listing 124: Finding Maximum Order Value

Output:**Output**

o_id	order_total
1	66600.00

(1 row)

Explanation: Order 1 (Arjun's first order) had Laptop + 2 Mice**Customers with Multiple Order Dates**

```

1 -- Customers with orders on different dates
2 SELECT c.c_name
3 FROM customers c
4 JOIN orders o ON c.c_id = o.c_id
5 GROUP BY c.c_name
6 HAVING COUNT(DISTINCT o.o_date) > 1;

```

Listing 125: Repeat Customers

Output:**Output**

c_name
Arjun

(1 row)

Note: Only Arjun placed orders on multiple dates**Popular Products (Ordered by 2+ Customers)**

```

1 -- Products ordered by more than 2 customers
2 SELECT p.p_name
3 FROM products p
4 JOIN order_items oi ON p.p_id = oi.p_id
5 JOIN orders o ON oi.o_id = o.o_id
6 GROUP BY p.p_name
7 HAVING COUNT(DISTINCT o.c_id) > 2;

```

Listing 126: Product Popularity

Output:

Output

```
p_name
-----
(0 rows)
```

Note: No product ordered by 3+ different customers in our dataset

Total Items per Order

```
-- Total items per order
SELECT o.o_id, o.o_date, SUM(oi.quantity) AS total_items
FROM orders o
JOIN order_items oi ON o.o_id = oi.o_id
GROUP BY o.o_id, o.o_date
ORDER BY o.o_id;
```

Listing 127: Order Summary

Output:

Output

o_id	o_date	total_items
1	2025-01-01	3
2	2025-01-02	1
3	2025-01-03	1
4	2025-01-04	1
5	2025-01-05	2

(5 rows)

Customers Who Ordered All Products

```
-- Customers who ordered all products
SELECT c.c_name
FROM customers c
JOIN orders o ON c.c_id = o.c_id
JOIN order_items oi ON o.o_id = oi.o_id
GROUP BY c.c_id, c.c_name
HAVING COUNT(DISTINCT oi.p_id) = (SELECT COUNT(*) FROM products
);
```

Listing 128: Complete Product Coverage

Output:

Output

```
c_name
-----
(0 rows)
```

Logic:

- Count distinct products ordered by each customer
- Compare to total product count
- Return customers where counts match

20 Views: Virtual Tables

20.1 Introduction

A view is a virtual table based on a SQL query. It doesn't store data but provides a way to simplify complex queries.

20.2 Why Use Views?

- **Simplification:** Hide complex JOIN logic
- **Security:** Restrict access to specific columns
- **Reusability:** Define once, use multiple times
- **Abstraction:** Present data in different formats

20.3 Creating Views

20.3.1 Syntax

Syntax

```

1 CREATE VIEW view_name AS
2 SELECT columns
3 FROM tables
4 WHERE conditions;
5
6 -- OR REPLACE to update existing view
7 CREATE OR REPLACE VIEW view_name AS
8 SELECT ...;
```

Billing Information View

```

1 -- Create view for billing information
2 CREATE OR REPLACE VIEW billing_info AS
3 SELECT
4     c.c_name,
5     o.o_id,
6     o.o_date,
7     p.p_name,
8     p.p_price,
9     oi.quantity,
10    SUM(p.p_price * oi.quantity) AS total_price
11   FROM customers c
12  JOIN orders o      ON c.c_id = o.c_id
13  JOIN order_items oi  ON o.o_id = oi.o_id
14  JOIN products p     ON p.p_id = oi.p_id
15 GROUP BY
16     c.c_name,
17     o.o_id,
18     o.o_date,
19     p.p_name,
20     p.p_price,
21     oi.quantity
22 ORDER BY
```

```
23     o.o_id;
```

Listing 129: Complex View Creation

Result: View created successfully

20.4 Querying Views

Using the View

```
1 -- Query view like a regular table
2 SELECT * FROM billing_info;
```

Listing 130: Querying a View

Output:

Output

c_name	o_id	o_date	p_name	p_pr.	qty	total_price
Arjun	1	2025-01-01	Laptop	65000	1	65000
Arjun	1	2025-01-01	Mouse	800	2	1600
Priya	2	2025-01-02	Keyboard	1500	1	1500
Arjun	3	2025-01-03	Headphones	2500	1	2500
Rahul	4	2025-01-04	Monitor	12000	1	12000
Sneha	5	2025-01-05	Mouse	800	1	800
Sneha	5	2025-01-05	Keyboard	1500	1	1500

(7 rows)

20.5 Aggregation on Views

Filtering View Results

```
1 -- Total revenue per product (filtered)
2 SELECT
3     p_name,
4     SUM(total_price) AS product_revenue
5 FROM billing_info
6 GROUP BY p_name
7 HAVING SUM(total_price) > 1500;
```

Listing 131: Product Revenue from View

Output:

Output

p_name		product_revenue
Laptop		65000.00
Mouse		2400.00
Headphones		2500.00
Keyboard		3000.00
Monitor		12000.00
(5 rows)		

Advantage: Complex JOIN hidden, simple query on view

20.6 ROLLUP: Grand Totals

20.6.1 Introduction

ROLLUP is a GROUP BY extension that creates subtotals and grand totals.

20.6.2 Syntax

Syntax

```

1 SELECT column, aggregate_function(column)
2 FROM table
3 GROUP BY ROLLUP (column);

```

Product Revenue with Grand Total

```

1 -- Product-wise totals + Grand Total
2 SELECT
3     COALESCE(p_name, 'Total') AS product_name,
4     SUM(total_price) AS total_revenue
5 FROM billing_info
6 GROUP BY ROLLUP (p_name)
7 ORDER BY total_revenue;

```

Listing 132: Using ROLLUP

Output:

Output

product_name		total_revenue
Keyboard		3000.00
Mouse		2400.00
Headphones		2500.00
Monitor		12000.00
Laptop		65000.00
Total		84900.00
(6 rows)		

Explanation:

- ROLLUP creates an extra row with NULL for grouped column
- COALESCE replaces NULL with 'Total'
- Grand total = sum of all product revenues

20.7 Managing Views

20.7.1 Dropping Views

Syntax

```
1 DROP VIEW view_name;
2
3 -- Drop if exists (no error if doesn't exist)
4 DROP VIEW IF EXISTS view_name;
```

20.7.2 Listing Views

```
-- PostgreSQL command to list views
\dv

-- Or using SQL
SELECT table_name
FROM information_schema.views
WHERE table_schema = 'public';
```

21 Stored Procedures and Functions

21.1 Introduction

Stored procedures and functions are reusable SQL code blocks stored in the database.

21.2 Stored Procedures

21.2.1 What is a Stored Procedure?

Definition

A **stored procedure** is a prepared SQL code that you can save and reuse. It can accept parameters and perform operations like INSERT, UPDATE, DELETE.

21.2.2 Syntax

Syntax

```

1 CREATE OR REPLACE PROCEDURE procedure_name(
2     parameter1 datatype,
3     parameter2 datatype
4 )
5 LANGUAGE plpgsql
6 AS $$*
7 BEGIN
8     -- SQL statements here
9 END;
10 $$;
11
12 -- Call procedure
13 CALL procedure_name(value1, value2);

```

Update Salary Procedure

```

1 -- Procedure to update employee salary
2 CREATE OR REPLACE PROCEDURE upd_salary(
3     p_sal NUMERIC(10,2),
4     p_id INT
5 )
6 LANGUAGE plpgsql
7 AS $$*
8 BEGIN
9     UPDATE employees
10    SET salary = p_sal
11   WHERE emp_id = p_id;
12 END;
13 $$;
14
15 -- Call the procedure
16 CALL upd_salary(100000, 10);

```

Listing 133: Simple UPDATE Procedure

Result: Employee 10's salary updated to 100,000

Insert Employee Procedure

```

1 -- Procedure to add new employee
2 CREATE OR REPLACE PROCEDURE add_employee(
3     fname VARCHAR,
4     dept VARCHAR,
5     salary NUMERIC,
6     age INT,
7     hire_date DATE,
8     bonus NUMERIC,
9     status VARCHAR
10 )
11 LANGUAGE plpgsql
12 AS $$*
13 BEGIN
14     INSERT INTO employees (
15         fname, dept, salary, age, hire_date, bonus, status
16     )
17     VALUES (
18         fname, dept, salary, age, hire_date, bonus, status
19     );
20 END;
21 $$*
22
23 -- Call the procedure
24 CALL add_employee(
25     'Manish', 'IT', 75000, 26, '2024-01-12', 5000, 'ACTIVE'
26 );

```

Listing 134: INSERT Procedure with Multiple Parameters

Result: New employee added to database

21.3 User-Defined Functions (UDF)

21.3.1 What is a UDF?

Definition

A **User-Defined Function** returns a value or table and can be used in SELECT statements, unlike procedures.

21.3.2 Syntax

Syntax

```

1 CREATE OR REPLACE FUNCTION function_name(
2     parameter1 datatype
3 )
4 RETURNS return_type
5 LANGUAGE plpgsql
6 AS $$*
7 BEGIN
8     -- Logic here
9     RETURN result;

```

```

10 END;
11 $$;
```

Highest Paid Employee Function

```

1 -- Function to find highest paid employee(s) in a department
2 CREATE OR REPLACE FUNCTION max_salary_dept(p_dept_name VARCHAR)
3 RETURNS TABLE (
4     emp_id INT,
5     fname VARCHAR,
6     salary NUMERIC
7 )
8 LANGUAGE plpgsql
9 AS $$
10 BEGIN
11     RETURN QUERY
12     SELECT e.emp_id, e.fname, e.salary
13     FROM employees e
14     WHERE e.dept = p_dept_name
15     AND e.salary = (
16         SELECT MAX(e2.salary)
17         FROM employees e2
18         WHERE e2.dept = p_dept_name
19     );
20 END;
21 $$;
22
23 -- Call the function
24 SELECT * FROM max_salary_dept('IT');
```

Listing 135: Function Returning Table

Output:

Output

emp_id	fname	salary
10	Vikram	75000.00

(1 row)

21.4 Procedures vs Functions

Aspect	Procedure	Function
Returns	Nothing (void)	Value or table
Call method	CALL proc()	SELECT func()
Use in SELECT	No	Yes
Purpose	Perform actions	Return data
Transactions	Can manage	Read-only

22 Window Functions: Advanced Analytics

22.1 Introduction

Window functions perform calculations across a set of rows related to the current row, without collapsing rows like GROUP BY.

22.2 Window Function Syntax

Syntax

```

1 function_name() OVER (
2     [PARTITION BY column]
3     [ORDER BY column]
4     [ROWS/RANGE specification]
5 )

```

22.3 Running Totals and Averages

Running Total of Salary

```

1 -- Running total of salary (company-wide)
2 SELECT
3     fname,
4     salary,
5     SUM(salary) OVER (ORDER BY salary) AS running_total
6 FROM employees;

```

Listing 136: Cumulative Sum

Output:

Output

fname	salary	running_total
Karan	39000.00	39000.00
Anita	43000.00	82000.00
Priya	45000.00	127000.00
Neha	48000.00	175000.00
Pooja	51000.00	226000.00
Amit	52000.00	278000.00
Raj	60000.00	338000.00

(7 rows)

Explanation:

- Each row shows cumulative sum up to that point
- Ordered by salary (ascending)
- No grouping - all rows preserved

Running Average

```

1 -- Running average of salary
2 SELECT
3     fname,
4     salary,
5     AVG(salary) OVER (ORDER BY salary) AS running_avg
6 FROM employees;

```

Listing 137: Moving Average

Output:

Output

fname	salary	running_avg
Karan	39000.00	39000.0000000000000000000000
Anita	43000.00	41000.0000000000000000000000
Priya	45000.00	42333.3333333333333333
Neha	48000.00	43750.0000000000000000000000
(4 rows)		

22.4 PARTITION BY: Department-wise Analysis

Highest Paid Employee per Department

```

1 -- Highest paid employee in each department
2 SELECT fname, dept, salary
3 FROM (
4     SELECT
5         fname,
6         dept,
7         salary,
8         RANK() OVER (PARTITION BY dept ORDER BY salary DESC) AS
9             rnk
10    FROM employees
11 ) t
11 WHERE rnk = 1;

```

Listing 138: Using RANK() with PARTITION BY

Output:

Output

fname	dept	salary
Vikram	IT	75000.00
Pooja	HR	51000.00
Suman	Finance	70000.00
Amit	Marketing	52000.00
(4 rows)		

How it works:

1. PARTITION BY dept: Separate ranking for each department
2. ORDER BY salary DESC: Rank by highest salary
3. RANK(): Assigns rank (1 = highest)
4. Filter to rank = 1

Employee Count per Department

```

1 -- Employee count per department (shown on each row)
2 SELECT DISTINCT
3     dept,
4     COUNT(*) OVER (PARTITION BY dept) AS dept_count
5 FROM employees;

```

Listing 139: COUNT() Window Function

Output:

Output

dept		dept_count
IT		4
HR		2
Finance		2
Marketing		2
(4 rows)		

22.5 Ranking Functions

Function	Description
ROW_NUMBER()	Unique sequential number (no ties)
RANK()	Ranking with gaps for ties
DENSE_RANK()	Ranking without gaps

Top 3 Highest Paid per Department

```

1 -- Top 3 highest paid employees per department
2 SELECT fname, dept, salary
3 FROM (
4     SELECT
5         fname,
6         dept,
7         salary,
8         RANK() OVER (PARTITION BY dept ORDER BY salary DESC) AS
9             rk
10    FROM employees
11 ) t
12 WHERE rk <= 3
13 ORDER BY dept, salary DESC;

```

Listing 140: Department Top 3

Output:**Output**

fname	dept	salary
Suman	Finance	70000.00
Ravi	Finance	62000.00
Pooja	HR	51000.00
Priya	HR	45000.00
Vikram	IT	75000.00
Raj	IT	60000.00
Amit	IT	52000.00
Neha	Marketing	48000.00
Anita	Marketing	43000.00

(9 rows)

22.6 Cumulative Sum per Department

Running Total by Department

```

1 -- Cumulative salary per department
2 SELECT
3     fname,
4     dept,
5     salary,
6     SUM(salary) OVER (
7         PARTITION BY dept
8         ORDER BY salary
9         ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW
10    ) AS running_total
11 FROM employees
12 ORDER BY dept, salary;

```

Listing 141: Department-wise Cumulative Salary

Output:**Output**

fname	dept	salary	running_total
Ravi	Finance	62000.00	62000.00
Suman	Finance	70000.00	132000.00
Priya	HR	45000.00	45000.00
Pooja	HR	51000.00	96000.00
Karan	IT	39000.00	39000.00
Amit	IT	52000.00	91000.00
Raj	IT	60000.00	151000.00
Vikram	IT	75000.00	226000.00

(8 rows)

Frame Specification:

ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW

- Includes all rows from start to current
- Creates cumulative effect

23 LEAD and LAG Functions

23.1 Introduction

LEAD and LAG access data from subsequent or previous rows without self-joins.

23.2 LAG Function

LAG Definition

`LAG(column, offset)` accesses data from **previous** rows.

Salary Difference from Previous Employee

```

1 -- Salary difference from previous employee in department
2 SELECT
3     fname,
4     dept,
5     salary,
6     salary - LAG(salary) OVER (
7         PARTITION BY dept ORDER BY salary
8     ) AS diff_from_prev
9 FROM employees;

```

Listing 142: Using LAG

Output:

Output

fname	dept	salary	diff_from_prev
Ravi	Finance	62000.00	NULL
Suman	Finance	70000.00	8000.00
Priya	HR	45000.00	NULL
Pooja	HR	51000.00	6000.00
Karan	IT	39000.00	NULL
Amit	IT	52000.00	13000.00
Raj	IT	60000.00	8000.00
Vikram	IT	75000.00	15000.00
(8 rows)			

Note: First row in each partition has NULL (no previous row)

Days Between Hires

```

1 -- Days gap between hires in department
2 SELECT
3     fname,
4     dept,
5     hire_date,
6     hire_date - LAG(hire_date) OVER (
7         PARTITION BY dept ORDER BY hire_date
8     ) AS days_gap

```

```
9 FROM employees;
```

Listing 143: Date Difference with LAG

Output:**Output**

fname	dept	hire_date	days_gap
Suman	Finance	2018-03-20	NULL
Ravi	Finance	2016-05-12	NULL
Pooja	HR	2017-09-30	NULL
Priya	HR	2020-01-10	832
(4 rows)			

23.3 LEAD Function

LEAD DefinitionLEAD(column, offset) accesses data from **next** rows.**Next Employee's Salary**

```
1 -- Next salary in department
2 SELECT
3     fname,
4     dept,
5     salary,
6     LEAD(salary) OVER (
7         PARTITION BY dept ORDER BY salary
8     ) AS next_salary
9 FROM employees;
```

Listing 144: Using LEAD

Output:**Output**

fname	dept	salary	next_salary
Ravi	Finance	62000.00	70000.00
Suman	Finance	70000.00	NULL
Priya	HR	45000.00	51000.00
Pooja	HR	51000.00	NULL
Karan	IT	39000.00	52000.00
Amit	IT	52000.00	60000.00
Raj	IT	60000.00	75000.00
Vikram	IT	75000.00	NULL
(8 rows)			

Note: Last row in each partition has NULL (no next row)

24 CTEs: Common Table Expressions

24.1 Introduction

CTEs create temporary named result sets that exist only during query execution. They improve readability and organization of complex queries.

24.2 Syntax

Syntax

```

1 WITH cte_name AS (
2     SELECT ...
3 )
4 SELECT * FROM cte_name;
5
6 -- Multiple CTEs
7 WITH
8     cte1 AS (SELECT ...),
9     cte2 AS (SELECT ...)
10 SELECT * FROM cte1 JOIN cte2;
```

24.3 CTE with GROUP BY

Employees Above Department Average (GROUP BY Method)

```

-- Employees earning above department average (GROUP BY CTE)
1 WITH avg_salary_dept AS (
2     SELECT dept, AVG(salary) AS avg_sal
3     FROM employees
4     GROUP BY dept
5 )
6 SELECT
7     e.emp_id,
8     e.fname,
9     e.dept,
10    e.salary,
11    a.avg_sal
12   a.avg_sal
13  FROM employees e
14  JOIN avg_salary_dept a
15  ON e.dept = a.dept
16 WHERE e.salary > a.avg_sal;
```

Listing 145: Using CTE with Aggregation

Output:

Output

emp_id	fname	dept	salary	avg_sal
4	Suman	Finance	70000.00	66000.000000000000000000
7	Pooja	HR	51000.00	48000.000000000000000000
1	Raj	IT	60000.00	56500.000000000000000000
10	Vikram	IT	75000.00	56500.000000000000000000
6	Amit	Marketing	52000.00	45500.000000000000000000

(5 rows)

How it works:

1. CTE calculates average salary per department
2. Main query joins employees with CTE
3. Filters employees earning above their department's average

24.4 CTE with Window Functions

Employees Above Department Average (Window Method)

```

1 -- Employees earning above department average (WINDOW + CTE)
2 WITH dept_avg AS (
3   SELECT
4     emp_id,
5     fname,
6     dept,
7     salary,
8     AVG(salary) OVER (PARTITION BY dept) AS avg_dept_salary
9   FROM employees
10 )
11 SELECT
12   emp_id,
13   fname,
14   dept,
15   salary,
16   ROUND(avg_dept_salary, 2) AS avg_dept_salary
17 FROM dept_avg
18 WHERE salary > avg_dept_salary;

```

Listing 146: CTE with Window Function

Output:

Output

emp_id	fname	dept	salary	avg_dept_salary
1	Raj	IT	60000.00	56500.00
4	Suman	Finance	70000.00	66000.00
6	Amit	Marketing	52000.00	45500.00
7	Pooja	HR	51000.00	48000.00
10	Vikram	IT	75000.00	56500.00

(5 rows)

Difference from GROUP BY Method:

- Window function doesn't collapse rows
- Average appears on every employee row
- No explicit JOIN needed
- More concise for this use case

24.5 Multiple CTEs

Chained CTEs for Complex Logic

```

1 -- Highest paid employees per department using multiple CTEs
2 WITH ranked_employees AS (
3     SELECT
4         emp_id,
5         fname,
6         dept,
7         salary,
8         RANK() OVER (PARTITION BY dept ORDER BY salary DESC) AS
9             rk
10    FROM employees
11 ), top_paid AS (
12     SELECT * FROM ranked_employees WHERE rk = 1
13 )
14 SELECT * FROM top_paid;

```

Listing 147: Multiple CTEs

Output:**Output**

emp_id	fname	dept	salary	rk
10	Vikram	IT	75000.00	1
7	Pooja	HR	51000.00	1
4	Suman	Finance	70000.00	1
6	Amit	Marketing	52000.00	1

(4 rows)

Execution Flow:

1. First CTE ranks employees within departments
2. Second CTE filters to top-ranked only
3. Final SELECT retrieves results

24.6 CTE vs Subquery

Aspect	CTE	Subquery
Readability	More readable	Can be nested, harder
Reusability	Can reference multiple times	Must repeat code
Debugging	Easier to test separately	Harder to isolate
Recursion	Supports recursive queries	Cannot recurse
Performance	Similar	Similar

24.7 Advanced CTE Examples

Highest Salary per Department - Two Methods

Method 1: GROUP BY + JOIN

```

1 -- Highest salary per department (GROUP BY + JOIN)
2 WITH max_salary_dept AS (
3     SELECT dept, MAX(salary) AS max_salary
4     FROM employees
5     GROUP BY dept
6 )
7 SELECT e.fname, e.dept, e.salary
8 FROM employees e
9 JOIN max_salary_dept a
10 ON e.dept = a.dept
11 WHERE e.salary = a.max_salary;

```

Listing 148: Traditional Approach

Method 2: Window Function

```

1 -- Highest salary per department (WINDOW FUNCTION)
2 WITH max_sal_dept AS (
3     SELECT
4         fname,
5         dept,
6         salary,
7         MAX(salary) OVER (PARTITION BY dept) AS max_sal
8     FROM employees
9 )
10 SELECT fname, dept, salary
11 FROM max_sal_dept
12 WHERE salary = max_sal;

```

Listing 149: Window Function Approach

Output (Both Methods):

Output

fname	dept	salary
Vikram	IT	75000.00
Pooja	HR	51000.00
Suman	Finance	70000.00
Amit	Marketing	52000.00

(4 rows)

25 Triggers: Automated Database Actions

25.1 Introduction

Triggers are special stored procedures that automatically execute when specific database events occur (INSERT, UPDATE, DELETE).

25.2 Trigger Components

1. **Trigger Function:** The logic to execute
2. **Trigger:** Associates function with table event

25.3 Creating a Trigger

25.3.1 Syntax

Syntax

```

1 -- Step 1: Create trigger function
2 CREATE OR REPLACE FUNCTION function_name()
3 RETURNS TRIGGER
4 LANGUAGE plpgsql
5 AS $$
6 BEGIN
7     -- Logic using NEW or OLD records
8     RETURN NEW; -- or OLD
9 END;
10 $$;
11
12 -- Step 2: Create trigger
13 CREATE TRIGGER trigger_name
14 BEFORE/AFTER INSERT/UPDATE/DELETE ON table_name
15 FOR EACH ROW
16 EXECUTE FUNCTION function_name();

```

25.4 NEW and OLD Records

Keyword	Description
NEW	Contains new values (INSERT, UPDATE)
OLD	Contains old values (UPDATE, DELETE)

25.5 Example: Salary Validation Trigger

Preventing Negative Salaries

```

1 -- Step 1: Create trigger function
2 CREATE OR REPLACE FUNCTION check_salary()
3 RETURNS TRIGGER
4 LANGUAGE plpgsql
5 AS $$
6 BEGIN
7     -- If updated salary is negative, reset it to 0
8     IF NEW.salary < 0 THEN

```

```

9      NEW.salary := 0;
10     END IF;
11
12     RETURN NEW;
13 END;
14 $$;
15
16 -- Step 2: Create trigger
17 CREATE TRIGGER before_upd_salary
18 BEFORE UPDATE ON employees
19 FOR EACH ROW
20 EXECUTE FUNCTION check_salary();

```

Listing 150: BEFORE UPDATE Trigger

How it works:

1. Trigger fires BEFORE any UPDATE on employees
2. Checks if NEW.salary is negative
3. If negative, sets it to 0
4. Returns modified NEW record
5. Database saves the corrected value

25.6 Removing Conflicting Constraints

```

1 -- Remove existing CHECK constraint (if any)
2 -- Allows trigger-based validation instead
3 ALTER TABLE employees
4 DROP CONSTRAINT IF EXISTS employees_salary_check;

```

Listing 151: Drop CHECK Constraint

25.7 Testing the Trigger

Trigger Test Cases

```

1 -- Test 1: Attempt to set negative salary (trigger will correct
   it to 0)
2 CALL upd_salary(-50000, 10);
3
4 -- Check result
5 SELECT emp_id, fname, salary
6 FROM employees
7 WHERE emp_id = 10;

```

Listing 152: Testing Negative Salary

Output:

Output

```
emp_id | fname | salary
-----+-----+-----
 10 | Vikram |   0.00
(1 row)
```

Result: Trigger intercepted negative value and set it to 0

```
1 -- Test 2: Valid salary update
2 CALL updt_salary(50000, 10);
3
4 -- Check result
5 SELECT emp_id, fname, salary
6 FROM employees
7 WHERE emp_id = 10;
```

Listing 153: Valid Salary Update

Output:**Output**

```
emp_id | fname | salary
-----+-----+-----
 10 | Vikram | 50000.00
(1 row)
```

Result: Valid update proceeds normally

25.8 Trigger Types

Type	Description
BEFORE	Executes before the operation, can modify data
AFTER	Executes after operation, for logging/auditing
INSTEAD OF	Replaces the operation (used with views)

25.9 Common Trigger Use Cases

- Data Validation:** Ensure data integrity before saving
- Auditing:** Log changes to sensitive data
- Automatic Calculations:** Update computed fields
- Referential Integrity:** Cascade operations
- Business Rules:** Enforce complex constraints

25.10 Managing Triggers

25.10.1 Listing Triggers

```
-- List triggers on a table
\d employees

-- Or using SQL
SELECT trigger_name
FROM information_schema.triggers
WHERE event_object_table = 'employees';
```

25.10.2 Dropping Triggers

Syntax

```
1 DROP TRIGGER trigger_name ON table_name;
2
3 -- Also drop the function if no longer needed
4 DROP FUNCTION function_name();
```

26 PostgreSQL Best Practices and Summary

26.1 Key Concepts Covered

1. Database Fundamentals

- What databases are and why we need them
- PostgreSQL architecture and features
- SQL categories (DDL, DML, DQL, DCL, TCL)

2. DDL Operations

- CREATE TABLE with constraints
- ALTER TABLE operations
- Data types and their usage

3. DML Operations

- INSERT, UPDATE, DELETE
- TRUNCATE vs DELETE

4. Querying Data

- SELECT with WHERE clause
- Comparison and logical operators
- LIKE pattern matching
- ORDER BY, LIMIT, OFFSET
- DISTINCT for unique values

5. Aggregate Functions

- COUNT, SUM, AVG, MAX, MIN
- GROUP BY for grouping data
- HAVING for filtering groups

6. String Functions

- CONCAT, SUBSTRING, REPLACE
- LENGTH, UPPER, LOWER
- TRIM, POSITION

7. Subqueries

- Scalar subqueries
- Correlated subqueries
- EXISTS and NOT EXISTS

8. CASE Statements

- Conditional logic in queries
- Categorization and transformations

9. JOINs

- INNER JOIN, LEFT JOIN
- Multi-table joins
- JOIN with GROUP BY and HAVING

10. Advanced Features

- Views for query simplification
- ROLLUP for grand totals
- Stored procedures and functions
- Window functions and analytics
- LEAD/LAG for row comparisons
- CTEs for query organization
- Triggers for automated actions

26.2 SQL Best Practices

Writing Production-Quality SQL

1. Always Use WHERE with UPDATE/DELETE

- Never run UPDATE or DELETE without WHERE
- Double-check condition before execution
- Test with SELECT first

2. Use Explicit Column Names

- Avoid SELECT * in production
- Specify exact columns needed
- Better performance and maintainability

3. Proper Indexing

- Index columns used in WHERE, JOIN, ORDER BY
- Don't over-index (slows INSERT/UPDATE)
- Analyze query performance

4. Use Constraints

- PRIMARY KEY for unique identifiers
- FOREIGN KEY for referential integrity
- CHECK constraints for data validation
- NOT NULL where appropriate

5. Transaction Management

- Use transactions for related operations
- COMMIT on success, ROLLBACK on error

- Keep transactions short

6. Naming Conventions

- Use descriptive, consistent names
- snake_case for tables and columns
- Prefix constraints meaningfully

7. Security

- Never trust user input
- Use parameterized queries
- Implement least privilege access
- Regular backups

26.3 Query Optimization Tips

1. Use EXPLAIN ANALYZE

- Understand query execution plans
- Identify bottlenecks
- Optimize slow queries

2. Avoid N+1 Query Problems

- Use JOINs instead of multiple queries
- Fetch related data in one query

3. Use Window Functions Wisely

- More efficient than self-joins
- Better than correlated subqueries
- But can be memory-intensive

4. Limit Result Sets

- Use LIMIT for large datasets
- Implement pagination properly
- Don't fetch unnecessary data

26.4 Common Pitfalls to Avoid

Warning

1. Implicit Type Conversions

- Can prevent index usage
- Be explicit with data types

2. Using OR Instead of IN

- IN is more readable
- Often better optimized

3. Not Understanding NULL

- NULL is not equal to anything (including NULL)
- Use IS NULL, not = NULL
- Handle NULLs with COALESCE

4. Forgetting to Analyze Performance

- Small datasets hide performance issues
- Test with realistic data volumes
- Monitor query performance

26.5 Essential PostgreSQL Commands

Quick Reference

```
-- Database Operations
\l                                -- List databases
\c database_name                  -- Connect to database
CREATE DATABASE db_name;          -- Create database
DROP DATABASE db_name;            -- Delete database

-- Table Operations
\dt                               -- List tables
\d table_name                      -- Describe table
\d+ table_name                     -- Detailed table info

-- User Operations
\du                               -- List users
CREATE USER username;             -- Create user
GRANT ALL ON table TO user;      -- Grant permissions

-- Query Helpers
\x                                -- Toggle expanded display
\timing                            -- Show query execution time
\! clear                           -- Clear screen

-- Exiting
\q                                -- Quit psql
```

26.6 Next Steps for Learning

1. Practice Regularly

- Solve SQL challenges online
- Build real projects

- Experiment with different datasets

2. Learn Advanced Topics

- Recursive CTEs
- Full-text search
- JSON/JSONB operations
- Performance tuning
- Replication and high availability

3. Explore PostgreSQL Ecosystem

- Extensions (PostGIS, pg_stat_statements)
- Backup and recovery strategies
- Connection pooling (PgBouncer)
- Monitoring tools

4. Real-World Application

- Integrate with Python (psycopg2, SQLAlchemy)
- Build REST APIs with FastAPI
- Create data pipelines
- Implement analytics dashboards

End of PostgreSQL Complete Mastery Guide

*"Data is the new oil, but like oil, it's valuable only when refined."
SQL is the refinement tool.*

Keep querying, keep learning!

Master PostgreSQL to unlock the full potential of your data.